

1 Problem 1: Description of the Problem, Subscribe a ML/NN Based Solution

Process:

Definition:

Example:

2 Problem 2: Prescribe a strategy to optimize the NN.

Process:

Definition:

Example:

3 Problem 3: Explain step by step inference for basic algorithms (MLP, CNN, GNN, attention mechanism) in terms of numpy or basic tensor operations.

Process:

Definition:

Example:

4 Problem 4: Be able to explain why such a solution might work or fail.

Process:

Definition:

Example:

5 L2 ML

Summary:

- What does it mean when a computer program learns?
- What is a neural network?
- What is something that a AI model CANNOT do right now?

6 L3 Neural Networks

Summary:

- What is a neural network (NN)?
- How is a GLM and a neural network related?
- What is a representation? Other names for this?
- How do NN make predictions?
- How do NN learn?

7 L4 Neural Network Engineering

Summary:

- What is a MLP?
- What is an inductive bias and why might they be useful?
- What is the difference between hyperparameters and parameters?
- How do we optimize all parameters in a model?

8 L5 Optimizing Hyperparameters

Summary:

- What strategies can help a NN converge when training?
- What hyperparameters does a NN architecture have?
- How can we optimize parameters without gradients?
- DL requires a lot of data, what can we do when data is scarce?

8.1 Learning representations of data

8.1.1 AI/ML/DL

Definition:

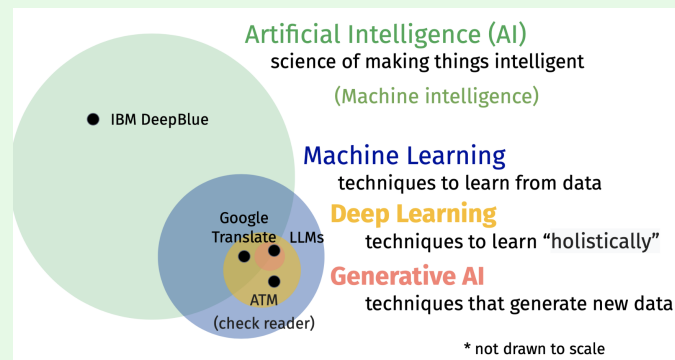


Figure 1:

8.1.2 Learning algorithms

Definition: “A computer program M is said to **learn** from **experience E** with respect to some class of **tasks T** and **performance measure P** , if its performance at **tasks in T** , as measured by **P** , improves with **experience E** ”

- **experience E** \sim Data

- **performance measure P** \sim Loss function, evaluation metric
- **tasks T** \sim “Prediction problem”
- **computer program M** \sim Model
- **learn** \sim Optimize

8.1.3 Linear models

Definition: Linear Regression

$$W \cdot x = y \quad (1)$$

- **E?** (x and y)
- **P?** mean squared error
- **T?** Predict y from x
- **M?** Linear model (W)
- **learn?** Analytical solution or gradient descent

Definition: Generalized Linear models in equations

$$\text{Link}(W \cdot x) = y \quad (2)$$

- x : Input features
- W : Linear transformation
- y : Output / target
- $\text{Link}(x)$: Warping function

Example:

1. If x has dim 50 and W projects to dimension 100, what is the shape of W ?
 - W is a 100×50 matrix
2. If W is learnable, how many parameters does W have?
 - $100 \times 50 = 5000$ parameters

Notes: How does a generalized linear model make a prediction? By either mapping to a line or separating data by a line (hyperplane)




Figure 2:

Notes: What can we do when the data cannot be separated by a line? Resort to different decision surfaces.

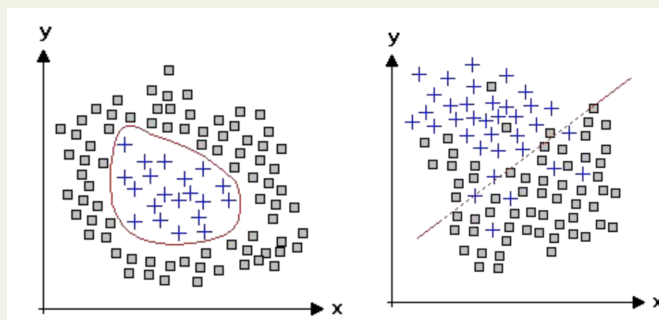


Figure 3:

8.1.4 Representations

Definition: Representation is a way of encoding data.

$$x \xrightarrow{\text{Representation}} z \quad (3)$$

- z : Feature vectors, embeddings, latent codes, intermediate activations, etc.

Notes:

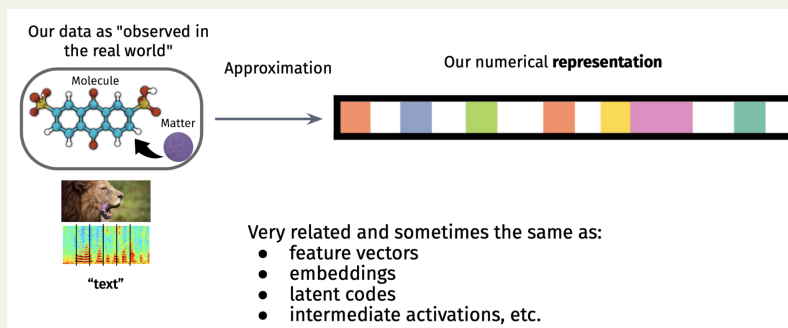


Figure 4:

8.2 Neural networks

Definition: Learnable (optimizable) transformations of data.

$$x \xrightarrow{\text{Model}} y \quad (4)$$

8.2.1 2-layer MLP

Definition: By stacking linear transforms with activation functions.

$$\text{Link}(W_2 \cdot \text{relu}(W_1 \cdot x)) = y \quad (5)$$

- x : Input features.
- W_1, W_2 : Linear transformations or Weight Matrices.
- $\text{relu}(x) = \max(0, x)$: Non-linear activation function, s.t. $f'(x) = 1$ if $x > 0$ and 0 otherwise.
- y : Output / target.
- **E?** (x and y)
- **P?** mean squared error

- **T?** Predict y from x
- **M?** Neural net (W_1, W_2)
- **learn?** gradient descent

Example:

1. What purpose does relu serve?
 - Introduces non-linearity into the model, allowing it to learn more complex functions.
2. If x has dim 50 and y dim 10, we have layer size of 50, how many parameters do we have?
 - W_1 is a 50×50 matrix, so it has $50 \times 50 = 2500$ parameters.
 - W_2 is a 10×50 matrix, so it has $10 \times 50 = 500$ parameters.
 - Total parameters: $2500 + 500 = 3000$ parameters. IS THIS CORRECT?

8.2.2 Geometric intuition

Notes: Decision surfaces Different ways of cutting up space to make predictions.

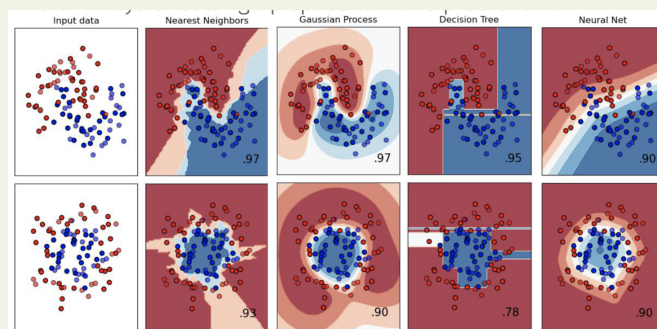


Figure 5:

Notes: Linear Transformation Transform data from one vector space to another

$$W \cdot x \quad (6)$$

Notes: SVD of Linear Transformation Factorizing matrices into geometrical transformations.

$$W = U \Sigma V^T \quad (7)$$

- U, V : Rotation
- Σ : Scaling

Notes: Affine Transformation

$$W \cdot x + b \quad (8)$$

- b : Bias vector
- Translate (b)
- Rotate (W -SVD)
- Reflect (W -SVD)
- Scale (W -SVD)
- Project up or down (dimensionality of $W\mathbf{x}$)

Notes: ReLU Rectified linear unit, which has a geometric effect of "gating", some info passes, some doesn't.




Figure 6:

Notes: Neural nets Learn to warp space to make better predictions.




Figure 7:




Figure 8:

8.2.3 Encoder-Decoder view

Definition:

$$x \xrightarrow{\text{Encoder}} z \xrightarrow{\text{Decoder}} y \quad (9)$$

- z : Embeddings, latent vectors, learned representations.

Example: Supervised Learning

$$x \xrightarrow{\text{Model}} y \quad (10)$$

- $\text{Model}(x) = \text{Decoder}(\text{Encoder}(x))$
- $\text{Decoder}(z) = \text{Pred}(z)$

Example: PCA

$$x \xrightarrow{\text{Encoder}} z \xrightarrow{\text{Decoder}} y \quad (11)$$

- $\text{Encoder}(x) = W \cdot x$
- $\text{Decoder}(z) = W^{-1} \cdot z$
- **E?** x
- **P?** Reconstruction loss
- **T?** Reduce dimension
- **M?** W
- **learn?** Eigendecompositions

Example: Neural Networks

$$x \xrightarrow{\text{Encoder}} z \xrightarrow{\text{Decoder}} y \quad (12)$$

- $\text{Encoder}(x) = \text{Neural Network}$
- $\text{Decoder}(z) = \text{Neural Network}$

8.2.4 Typical ML Pipeline**Notes:**

- Setup data (x, y)
- Define a model: $y = f(x, \theta)$
- Training algorithm to find θ
- Evaluate the model.

8.3 Recap**Summary:**

- What is a MLP? Vector-in vector-out optimizable, learnable transformation of data.
- What is an inductive bias and why might they be useful? Set of assumptions that the learner puts on a model for a task, makes an algorithm learn one pattern over another.
 - Let certain patterns be learnable (restricting the hypothesis space).
 - Last layer in the GLM: Restricting output values to 0 and 1.
- What is the difference between hyperparameters and parameters? Hyperparameters are set before training (usually discrete), parameters are learned during training (continuous to learn).
- How do we optimize all parameters in a model? Back propagation.

8.4 Optimizing Hyperparameters

Summary: