

ECE353 Lectures

Hanhee Lee

January 9, 2025

Contents

Definition:

Process:

Motivation:

Derivation:

Warning:

Summary:

Algorithm:

Example:

FAQ:

1 Review

1.1 Converting Between Binary, Hexadecimal, and Decimal

Process:

1. **Binary to Decimal:**
 - (a) Write down the binary number.
 - (b) Assign place values, starting from 2^0 on the rightmost digit.
 - (c) Multiply each binary digit by its corresponding power of 2.
 - (d) Add all the results together to get the decimal equivalent.
2. **Decimal to Binary:**
 - (a) Divide the decimal number by 2.
 - (b) Record the remainder (0 or 1).
 - (c) Repeat the division process with the quotient until the quotient is 0.
 - (d) Write the remainders in reverse order to obtain the binary equivalent.
3. **Binary to Hexadecimal:**
 - (a) Group the binary number into groups of 4 digits, starting from the right. Add leading zeros if necessary.
 - (b) Convert each 4-digit binary group to its hexadecimal equivalent using the binary-to-hex mapping (e.g., 0000 = 0, 0001 = 1, 1110 = E).
 - (c) Combine the hexadecimal digits to get the hexadecimal equivalent.
4. **Hexadecimal to Binary:**
 - (a) Write down each hexadecimal digit.
 - (b) Replace each hexadecimal digit with its 4-bit binary equivalent.
 - (c) Combine the binary groups to get the binary equivalent.
5. **Decimal to Hexadecimal:**
 - (a) Divide the decimal number by 16.
 - (b) Record the remainder as a hexadecimal digit (0–9 or A–F).
 - (c) Repeat the division process with the quotient until the quotient is 0.
 - (d) Write the remainders in reverse order to obtain the hexadecimal equivalent.
6. **Hexadecimal to Decimal:**
 - (a) Write down the hexadecimal number.
 - (b) Assign place values, starting from 16^0 on the rightmost digit.
 - (c) Multiply each hexadecimal digit by its corresponding power of 16, converting any letters (A–F) to decimal values (A=10, B=11, etc.).
 - (d) Add all the results together to get the decimal equivalent.

1.2 Little-endian and Big-endian

Definition:

- **Little-endian:** In the little-endian format, the least significant byte (LSB) of a multi-byte data value is stored at the lowest memory address, and the most significant byte (MSB) is stored at the highest memory address.
- **Big-endian:** In the big-endian format, the most significant byte (MSB) of a multi-byte data value is stored at the lowest memory address, and the least significant byte (LSB) is stored at the highest memory address.

Example:

- For example, the hexadecimal value 0x12345678 would be stored in memory as:

78 56 34 12

- For example, the hexadecimal value 0x12345678 would be stored in memory as:

12 34 56 78

1.3 Memory

Summary: Table, int*, &a, int**a, *a, int[5], etc.

2 Why Systems Software?

Summary:

2.1 Useful Terminal Commands

Summary:

- `./hello-world-linux-aarch64` to run hello world.
- `readelf -a <FILE>` to see the ELF header.
- `strace <PROGRAM>` to trace all the system calls a process makes on Linux.

2.2 Three OS Concepts

Definition:

1. **Virtualization:** Share one resource by mimicking multiple independent copies.
2. **Concurrency:** Handle multiple things happening at the same time.
3. **Persistence:** Retain data consistency even without power.

2.3 OS Manages Resources

Definition: Insert picture.

2.4 Program

Definition: A file containing all the instructions and data required to run.

2.5 Process:

Definition: An instance of running a program.




Figure 1: Process

2.5.1 Basic Requirements for a Process

Definition: Insert picture w/ virtual memory.

2.6 Process (Abstraction)

2.6.1 Static

Definition: Only able to use the global variable in the current C file.

2.6.2 Motivation for Virtualization

Motivation: How to run two different programs at the same time? Insert code.

- Was the address of local the same b/w 2 processes? Different address in physical memory b/w different processes.
- Was the address of global the same b/w 2 processes? Same address in physical memory b/w different processes, but uses virtual memory.
- What else may be needed for a process?

Warning: Local variables are stored on the stack.

2.6.3 Does the OS allocate different stacks for each process?

Definition: The stacks for each process need to be in physical memory. One option is the operating system just allocates any unused memory for the stack.

-

2.6.4 What about global variables?

Definition: The compiler needs to pick an address (random) for each variable when you compile.

- What if we had a global registry of addresses? Impossible (too much space and know memory addresses ahead of time).

3 Kernels

Summary:

- The kernel is the part of the operating system (OS) that interacts with hardware (it runs in kernel mode).
- System calls are the interface between user and kernel mode:
 - Every program must use this interface!
- File format and instructions to define a simple “Hello world” (in 168 bytes):
 - Difference between API and ABI.
 - How to explore system calls.
- Different kernel architectures shift how much code runs in kernel mode.

FAQ:

- What is difference b/w printf and write?

3.1 File Descriptor (Abstraction)

Motivation: Since our processes are independent, we need an explicit way to transfer data.

Definition:

1. **IPC:** Inter-process communication is transferring data b/w two processes.
2. **File Descriptor:** A resource that users may either read bytes from or write bytes to (identified by an index stored in a process).
 - e.g. File or terminal.
 - e.g. 0 is standard input, 1 is standard output, and 2 is standard error.

3.2 System Calls

Definition: System calls are the interface b/w user and kernel mode.

3.2.1 System Calls Make Requests to the Operating System

Definition:

```
1 ssize_t write(int fd, const void *buf, size_t count);
```

- Description: writes bytes from a byte array to a file descriptor
 - fd: the file descriptor
 - buf: the address of the start of the byte array (called a buffer)
 - count: how many bytes to write from the buffer

```
1 void exit_group(int status);
```

- Description: exits the current process and sets an exit status code
 - status: the exit status code (0–255)

Example: Hypothetical "Hello World" Program

```
1 void _start(void) {
2     write(1, "Hello world\n", 12);
3     exit_group(0);
4 }
```

Warning: System calls uses registers, while C is stack based.

3.3 API Tells You What and ABI Tells You How

Definition:

- Application Programming Interface (API) abstracts the details and describes the arguments and return value of a function.
- Application Binary Interface (ABI) specifies the details, specifically how to pass arguments and where the return value is.

3.4 Magic

Definition: The "magic bytes" refer to the first 4 bytes of a file that uniquely identify the file format.

3.4.1 Programs on Linux Use the ELF File Format

Definition: Executable and Linkable Format (ELF) specifies both executables and libraries.

- Always starts with the 4 bytes: 0x7F 0x45 0x4C 0x46 or with ASCII encoding: DEL 'E' 'L' 'F'

Example: Hello World ELF File

1. 168 Byte Program:

- Tells the OS to load the entire executable file into memory at address 0x10000.
- The file header is 64 bytes, and the “program header” is 56 bytes (120 bytes total).
- The next 36 bytes are instructions, followed by 12 bytes for the string:
 - "Hello world\n"
 - Instructions start at 0x10078 (0x78 is 120).
 - The string (data) starts at 0x1009C (0x9C is 156).



Figure 2: ELF File Division

2. **C Program:** Takes 500 bytes.
3. **Python Program:** Takes 2000 bytes.
4. **Java Program:** Takes 2000000 bytes.

3.5 Kernel

Definition: Kernel is a core part of the operating system that interacts with hardware that runs in kernel mode.

3.5.1 Kernel Mode

Definition: Kernel mode is a privilege level on your CPU that gives access to more instructions.

3.5.2 Levels of Privelege

Definition:

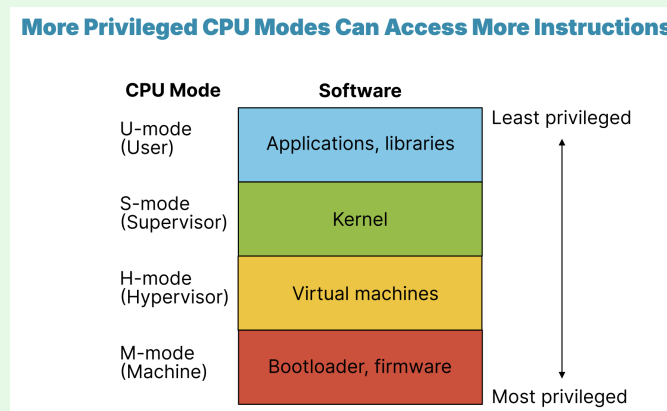


Figure 3: Levels of Privelege

3.5.3 System Calls Transition Between User and Kernel Mode

Definition:

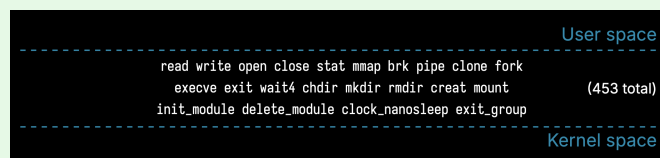


Figure 4: System Calls Transition

3.5.4 Different Tpyes of Kernel Architectures

Definition:

- **Monolithic Kernel:** All the services are in the kernel.




Figure 5: Monolithic Kernel

- **Microkernel:** Only the essential services are in the kernel.




Figure 6: Microkernel

- **Hybrid Kernel:** A mix of monolithic and microkernel.
- **Nanokernel and picokernel:** Even smaller services than microkernel.

Warning: Short answer question.

4 Libraries

4.1 What is an Operating System?

Definition: An operating system consists of a kernel and libraries required for your application.

Linux distributions may be considered GNU/Linux, where GNU distributes the standard C library and common utilities.

Warning: OS have different libraries for different applications.

4.2 Normal Compilation in C

4.3 Static Libraries and Dynamic Libraries

4.3.1 Comparison of Static and Dynamic Libraries

