ECE368 Cheatsheet

Hanhee Lee

January 26, 2025

Contents

Summary: On second thoughts, the lecture notes he posts are good, so I think I'll just do the cheatsheet.

W1 (LG-IPPR 1.1, 1.2; Murphy $2.1 - 2.3$)

1   L1: Probability Review

FAQ:
- How to study? Practice, practice.
- What textbooks? Use 2024 version of Murphy, Leon Garcia as main reference, Bishop, 4th textbook is intro.
- How is HW graded? Effort, and tutorials are used to explain soln.

1.1   Sample Space

Motivation: If you have 4 sheeps and a flea, the probability that starting from sheep 1, the flea will jump to sheep 4 in 10 steps is 0.2.
- Ambigious as there are 2 different interpretations for the sample space (i.e. space of probability is not clear):
  - Set of sheeps
  - Set of number of steps

1.2 Probability Definitions

Definition:
- **Random Experiment:** An outcome (realization) for each run.
- **Sample Space** Ω**:** Set of all possible outcomes.
- **Events:** (measurable) subsets of Ω.
- **Probability of Event A:** $P[A] \equiv P[$'outcome is in A'$]$.

Example: Roll Fair Die
- $\Omega = \{1, 2, 3, 4, 5, 6\}$.
- $P[$'even number'$] = \dfrac{1}{2}$.

1.3 Axioms of Probability

Definition:
1. $P[A] \geq 0$ for all $A \in \Omega$.
2. $P[\Omega] = 1$.
3. If $A \cap B = \emptyset$, then $P[A \cup B] = P[A] + P[B]$ for all $A, B \in \Omega$.

Figure 1: 3rd Axiom

1.4 Conditional Probability

Definition:

$$P[A|B] = \frac{P[A \cap B]}{P[B]} \tag{1}$$

- $|$: Given event (data/obs.).
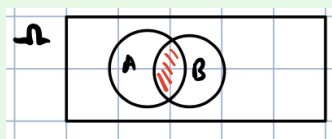
Figure 2: Conditional Probability

Notes:
- Changing sample space to B.
- Conditional probability satisfy the 3 axioms (i.e. are probabilities), can be viewed as probability measure on new sample space B.

1.4.1 Consequences of Conditional Probability

Definition:

$$P[A \cap B] = P[A|B]P[B] = P[B|A]P[A] \tag{2}$$

1.4.2   Independence

> **Definition**: A and B are independent iff
>
> $$P[A \cap B] = P[A]P[B] \iff P[A|B] = P[A] \iff P[B|A] = P[B] \tag{3}$$

1.4.3   Importance of Labelling

Example: **Toss 2 Fair Coins**
1. **Given:** Given that one of the coins is heads, what is the probability that the other coin is tails?
2. **Wrong Solution:** $\frac{1}{2}$ since $\{HH, HT, TH, TT\}$, so $P[T|H] = \frac{1}{2}$, which assumes that the coins are distinguishable (i.e. coin #1 is heads)
3. **Correct Solution:** $\frac{2}{3}$ since $\{HH, HT, TH\}$ as we didn't specify which coin was heads, so $P[T|H] = \frac{2}{3}$, which assumes that the coins are indistinguishable.

2   L2: Probability Review

FAQ:

2.1   Total Probability

Definition: If $H_1, \ldots, H_n$ form a partition of Ω, then

$$P[A] = \sum_{i=1}^{n} P[A|H_i]P[H_i] \tag{4}$$
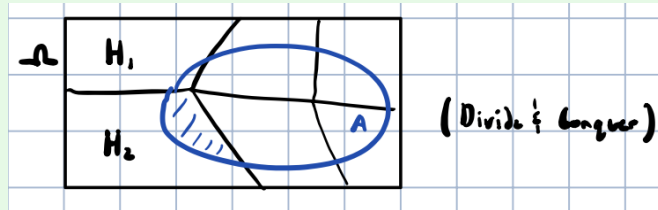
Figure 3: Total Probability

2.2   Bayes' Rule

Definition:

$$P[H_k|A] = \frac{P[H_k \cap A]}{P[A]} = \frac{P[A|H_k]P[H_k]}{\sum_{i=1}^{n} P[A|H_i]P[H_i]} \tag{5}$$

2.2.1   Posteriori Probability, Priori Probability (Prior), Likelihood

Definition:
- **Posteriori:** $P[H_k|A]$.
- **Priori:** $P[H_k]$.
- **Likelihood:** $P[A|H_k]$.

Example: Suppose a lie detector is 95% accurate, i.e. $P['\text{out=truth}'|'\text{in=truth}'] = 0.95$ and $P['\text{out=lie}'|'\text{in=lie}'] = 0.95$. It says that Mr. Ernst is lying. What is the probability Mr. Ernst is actually lying.
- **Observation:** $A = '\text{out=lie}'$.
- **Hypothesis:** $H_0 = '\text{in=lie}'$ and $H_1 = '\text{in=truth}'$.
- **Solution:** $P[H_0|A] = \dfrac{P[A|H_0]P[H_0]}{P[A|H_0]P[H_0] + P[A|H_1]P[H_1]} = \dfrac{0.95 \times P[H_0]}{0.95 \times P[H_0] + 0.05 \times (1 - P[H_0])}$.
- $H_0 = 0.01$: i.e. 1% of the population are liars, then $P[H_0|A] = \dfrac{0.95 \times 0.01}{0.95 \times 0.01 + 0.05 \times 0.99} = 0.16$.

Warning: Need to know priori probability.

2.2.2   Interpretation of Bayes' Rule

Notes: Taking one component of the total probability and normalizing it by the sum of all components.

2.3   Random Variables

Motivation: Coin Toss Mapping of each outcome to a real number
   • $w \in \Omega$ is the outcome of a coin toss, and X is the RV, so $H \to 0$ and $T \to 1$.
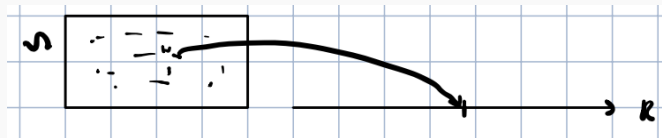
Figure 4: Random Variables

   • Mapping is deterministic function. RV is not random or variable.

Definition: Mapping from Ω to $\mathbb{R}$.

2.4   Distribution of RV

2.4.1   Cumulative Distribution Function (CDF) of RV

Definition:
$$F_X(x) \equiv P[X \leq x] \tag{6}$$

2.4.2   Discrete RV Probability Mass Function (PMF)

Definition:
$$P_X(x_j) \equiv P[X = x_j] \quad j = 1, 2, 3, \ldots \tag{7}$$

Example: Binonmial RV w/ (n, p)
$$P_X(k) = \binom{n}{k} p^k (1 - p)^{n-k} \tag{8}$$
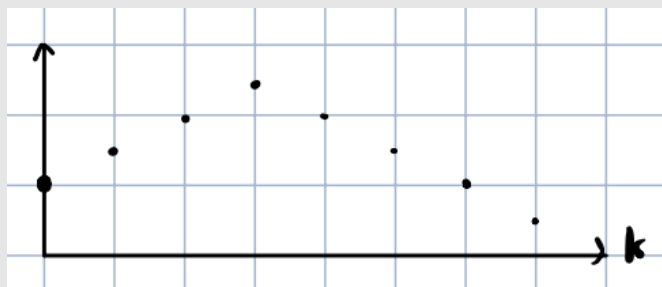
Figure 5: Binomial RV

2.4.3   Continuous RV Probability Density Function (PDF)

Definition:
$$f_X(x) \equiv \frac{d}{dx} F_X(x) \tag{9}$$
$$P[x < X < x + dx] = f_X(x)dx \tag{10}$$

Example: Gaussian RV w/ (μ, σ^2)

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \tag{11}$$
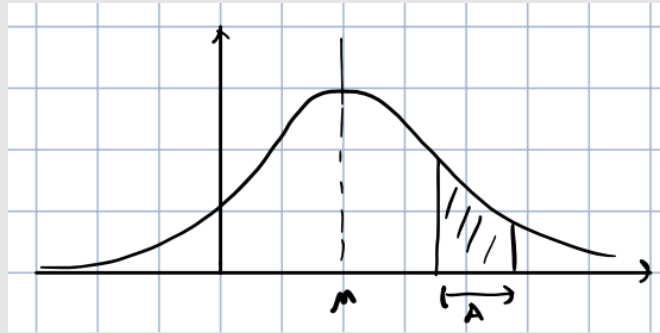
Figure 6: Gaussian RV

- $P[X \in A] = \int_A f_X(x)dx.$

Notes: Discrete RV has pdf w/ δ functions.

2.4.4   Conditional PMF/PDF

Definition:

$$P_X(x|A) \tag{12}$$

$$f_X(x|A) \tag{13}$$

Example: Continuous

$$f(x|X > a) = \begin{cases} \dfrac{f_X(x)}{P[X > a]} & \text{if } x > a \\ 0 & \text{otherwise} \end{cases} \tag{14}$$

Example: Geometric RV Geometric RV X w/ success probability p

$$P_X(k) = (1-p)^{k-1}p \tag{15}$$

- **Memoryless Property:** $P_X[k|X > m] = \dfrac{p(1-p)^{k-1}}{(1-p)^m} = p(1-p)^{k-m-1}.$
  - So it only cares about the additional trials (i.e. same as resetting after m trials).

2.5   Expected Values

Definition:

$$E[X] = \int_{-\infty}^{\infty} x f_X(x)dx \overset{\text{If int. values}}{=} \sum_{k=-\infty}^{\infty} k f_X(k) \tag{16}$$

$$E[h(X)] = \int_{-\infty}^{\infty} h(x) f_X(x)dx \overset{\text{If int. values}}{=} \sum_{k=-\infty}^{\infty} h(k) f_X(k) \tag{17}$$

$$\text{Var}[X] = E[(X - E[X])^2] = E[X^2] - E[X]^2 \tag{18}$$

$$E[X|A] = \int_{-\infty}^{\infty} x f_X(x|A) dx \tag{19}$$

Example: **Lottery Ticket (Geometric RV)**
1. **Given:** Buying one lottery ticket per week
    - Each ticket has $10^{-7} = p$ chance of winning the jackpot.
    - $X = $ '# of weeks to win jackpot'.
2. **Problem:** What is the expected number of weeks to win the jackpot?
3. **Solution:** $E[X] = \sum_{k=1}^{\infty} k(1-p)^{k-1}p = \ldots = \dfrac{1}{p} = 10^7$ weeks.
4. **Extension (Memoryless Property):** If I have already played for 999999 weeks, what is the expected number of weeks to win the jackpot? $E[X - 999999|X > 999999] = E[X] = 10^7$ weeks.

3   L3: Probability Review

> **FAQ**:

3.1   2 RVs

> **Notes**: RVs are neither random nor a variable.
>
> $$\underline{Z} = (X, Y)$$
>
>
>
> Figure 7: Mapping of RVs

3.2   Joint PMF/PDF

> **Definition**:
> $$P_{X,Y}(x, y) = P[X = x, Y = y] \tag{20}$$
>
> $$f_{X,Y}(x, y) = \frac{\partial^2}{\partial x \partial y} F_{X,Y}(x, y) \tag{21}$$
>
> $$P[(X, Y) \in A] = \int\!\!\int_{(x,y) \in A} f_{X,Y}(x, y)\, dx\, dy \tag{22}$$

> **Example**: Jointly Gaussian RVs X and Y with $(\mu_1, \mu_2, \sigma_1^2, \sigma_2^2, \rho)$
>
> $$f_{X,Y}(x, y) = \frac{1}{2\pi\sigma_1\sigma_2\sqrt{1 - \rho^2}} \exp\left\{ -\frac{1}{2(1 - \rho^2)} \left[\left(\frac{x - \mu_1}{\sigma_1}\right)^2 - 2\rho \left(\frac{x - \mu_1}{\sigma_1}\right) \left(\frac{y - \mu_2}{\sigma_2}\right) + \left(\frac{y - \mu_2}{\sigma_2}\right)^2 \right] \right\}$$

3.3   Expectations

> **Definition**:
> $$E[g(X, Y)] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} g(x, y) f_{X,Y}(x, y)\, dx\, dy$$

> **Notes**:
> - $g(X, Y)$ is also an RV, but inside the integral or sum, you use x and y as dummy variables to vary through the values of the RVs.

3.3.1   Correlation

> **Definition**:
> $$E[XY] \tag{23}$$

3.3.2 Covariance

Definition:
$$\text{Cov}[X, Y] = E[(X - \mu_X)(Y - \mu_Y)] = E[XY] - \mu_X \mu_Y = E[XY] - E[X]E[Y] \tag{24}$$

Notes:
- Mean shifted to 0.

3.3.3 Correlation Coefficient

Definition:
$$\rho_{X,Y} = E\left[\left(\frac{X - \mu_X}{\sigma_X}\right)\left(\frac{Y - \mu_Y}{\sigma_Y}\right)\right] = \frac{\text{Cov}[X, Y]}{\sigma_X \sigma_Y} \tag{25}$$

- $|\rho_{X,Y}| \leq 1$

Notes:
- Mean shifted to 0 and normalized by the standard deviation.

3.4 Marginal PMF/PDF

Definition:
$$P_X(x) = \sum_{j=1}^{\infty} P_{X,Y}(x, y_j), \quad P_Y(y) = \sum_{j=1}^{\infty} P_{X,Y}(x_j, y) \tag{26}$$

$$f_X(x) = \int_{-\infty}^{\infty} f_{X,Y}(x, y) \, dy, \quad f_Y(y) = \int_{-\infty}^{\infty} f_{X,Y}(x, y) \, dx \tag{27}$$

Notes:
- Total probability theorem is being used here.

Example: Jointly Gaussian X and Y:

$$f_X(x) = \int_{-\infty}^{\infty} f_{X,Y}(x, y) \, dy$$
$$= \dots \quad \text{(completing the square)}$$
$$= \frac{1}{\sqrt{2\pi}\sigma_1} e^{-\frac{(x - \mu_1)^2}{2\sigma_1^2}}, \quad \text{marginally Gaussian}$$

- Gaussian RVs has a property that the PDF of a single variable is equal to the marginal Gaussian of two variables.

3.5 Conditional PMF/PDF

Definition:
$$P_{X|Y}(x|y) \triangleq P[X = x | Y = y] = \frac{P_{X,Y}(x, y)}{P_Y(y)} \tag{28}$$

$$f_{X|Y}(x|y) \triangleq \frac{f_{X,Y}(x, y)}{f_Y(y)} \tag{29}$$

3.6   Bayes' Rule

Definition:
$$P_{Y|X}(x|y) = \frac{P_{X,Y}(x,y)}{P_X(x)} = \frac{P_{X|Y}(x|y)P_Y(y)}{\sum_{j=1}^{\infty} P_{X,Y}(x,y_j)P_Y(y_j)} \tag{30}$$

$$f_{Y|X}(y|x) = \frac{f_{X,Y}(x,y)}{f_X(x)} = \frac{f_{X|Y}(x|y)f_Y(y)}{\int_{-\infty}^{\infty} f_{X|Y}(x|y')f_Y(y')\,dy'} \tag{31}$$

3.7   Independent vs. Uncorrelated vs. Orthogonal

Definition:
1. Independent:
$$f_{X|Y}(x|y) = f_X(x)\ \forall y \Leftrightarrow f_{X,Y}(x,y) = f_X(x)f_Y(y) \tag{32}$$

2. Uncorrelated:
$$\mathrm{Cov}[X,Y] = 0 \quad \Leftrightarrow \quad \rho_{X,Y} = 0 \tag{33}$$

3. Orthogonal:
$$E[XY] = 0 \tag{34}$$

Theorem: If independent, then uncorrelated.

Derivation:
$$\text{Independent} \implies E[XY] = \int_{-\infty}^{\infty}\int_{-\infty}^{\infty} xy f_{X,Y}(x,y)\,dx\,dy$$
$$= \int_{-\infty}^{\infty}\int_{-\infty}^{\infty} xy f_X(x)f_Y(y)\,dx\,dy$$
$$= \left(\int_{-\infty}^{\infty} x f_X(x)\,dx\right)\left(\int_{-\infty}^{\infty} y f_Y(y)\,dy\right)$$
$$\implies E[XY] = E[X]E[Y]$$
$$\implies \mathrm{Cov}[X,Y] = 0, \quad \text{uncorrelated}$$
$$\nLeftarrow \text{ in general.}$$

Example: Jointly Gaussian RVs X and Y: If uncorrelated, i.e. $\rho_{X,Y} = 0$, then X and Y are independent.
$$f_{X,Y}(x,y) = \frac{1}{2\pi\sigma_1\sigma_2}\exp\left\{-\frac{1}{2}\left[\left(\frac{x-\mu_1}{\sigma_1}\right)^2 + \left(\frac{y-\mu_2}{\sigma_2}\right)^2\right]\right\}$$
$$= \frac{1}{\sqrt{2\pi}\sigma_1}e^{-\frac{(x-\mu_1)^2}{2\sigma_1^2}} \cdot \frac{1}{\sqrt{2\pi}\sigma_2}e^{-\frac{(y-\mu_2)^2}{2\sigma_2^2}}$$
$$= f_X(x)f_Y(y) \quad \text{independent}$$

3.8   Conditional Expectation

Definition:
$$E[Y] = E[E[Y|X]] \tag{35}$$
$$E[h(Y)] = E[E[h(Y)|X]] \tag{36}$$

Notes:
- $E[E[Y|X]]$ is w.r.t. X.
- $E[Y|X]$ is w.r.t. Y.

Derivation:

$$E[Y] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} y f_{X,Y}(x,y)\, dx\, dy$$

$$= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} y f_{Y|X}(y|x) f_X(x)\, dx\, dy$$

$$= \int_{-\infty}^{\infty} \left(\int_{-\infty}^{\infty} y f_{Y|X}(y|x)\, dy \right) f_X(x)\, dx$$

$$= \int_{-\infty}^{\infty} E[Y|X=x] f_X(x)\, dx \quad \text{(using the total probability theorem)}$$

$$= \int_{-\infty}^{\infty} g(x) f_X(x)\, dx$$

$$= E[g(X)]$$

$$= E[E[Y|X]].$$

Example:
1. **Given:** An unknown voltage. $X \sim \text{Uniform}(0,1)$. Measurement from a (bad) voltmeter: $Y \sim \text{Uniform}(0,X)$.

$$f_X(x) = \begin{cases} 1, & 0 < x < 1 \\ 0, & \text{otherwise} \end{cases}$$

$$f_{Y|X}(y|x) = \begin{cases} \dfrac{1}{x}, & 0 < y < x \\ 0, & \text{otherwise} \end{cases}$$

- **Note:** Area under PDF is 1.
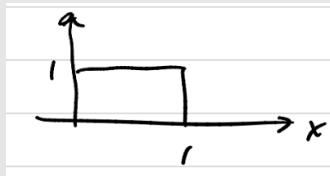
Figure 8: Uniform Distribution of X
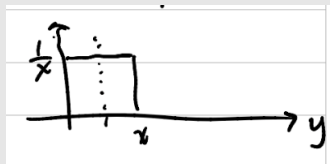
Figure 9: Uniform Distribution of Y

2. **Expected Value (Average Reading of Bad Voltmeter):**

$$E[Y] = E[E[Y|X]]$$
$$= E\left[\frac{X}{2}\right] \quad \text{Since in the middle of 0 and x}$$
$$= \frac{1}{2} \cdot E[X]$$
$$= \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{4} \quad \text{Since } E[X] \text{ (i.e. mean) is 0.5}$$

3. **The Long Way:**

$$f_Y(y) = \int_{-\infty}^{\infty} f_{Y|X}(y|x) f_X(x)\, dx$$
$$= \int_y^1 f_{Y|X}(y|x) f_X(x)\, dx$$
$$= \int_y^1 \frac{1}{x} \cdot 1\, dx$$
$$= -\ln y.$$
$$E[Y] = \int_0^1 y \cdot (-\ln y)\, dy = \cdots = \frac{1}{4}$$

4. **Question:** Suppose $Y = \frac{1}{8}$. What is "best" given X? This will be the quesiton for the rest of the course.

4   L4: Estimation of Sample Mean

FAQ:

4.1   Parameter Estimation:

Motivation: The readout of a sensor is $X = \theta + N$ volts

Figure 10:

- There is some noise N in the sensor, so we want to estimate the true value of θ (unknown parameter to be estimated)
    - e.g. Mean and/or variance of X.

4.2   Estimator:

Definition: Perform n independent and identically distributed (i.i.d.) measurements/observations of X: $X_1, X_2, \ldots, X_n$.

$$\hat{\Theta} = \hat{\Theta}(\underline{X}) = g(X_1, X_2, \ldots, X_n) \tag{37}$$

Figure 11:

4.2.1   Estimation Error:

Definition:

$$\hat{\Theta}(\underline{X}) - \theta \tag{38}$$

4.2.2   Unbiased

Definition: The estimator $\hat{\Theta}$ is unbiased if

$$\mathbb{E}[\hat{\Theta}(\underline{X})] = \theta \tag{39}$$

- **Asymptotically Unbiased:** $\lim_{n\to\infty} \mathbb{E}[\hat{\Theta}(\underline{X})] = \theta$ (big data)

4.2.3   Consistent

Definition: The estimator $\hat{\Theta}$ is consistent if $\hat{\Theta}(\underline{X}) \to \theta$ as $n \to \infty$, in probability, i.e., $\forall \epsilon > 0$,

$$\lim_{n\to\infty} P(|\hat{\Theta}(\underline{X}) - \theta| < \epsilon) \to 1 \tag{40}$$

as $n \to \infty$.

4.3   Sample Mean & Law of Large Numbers

Definition: Given a sequence of i.i.d. random variables (RVs), $X_1, X_2, \ldots, X_n$, w/ unknown mean μ, estimate μ. Let $S_n = X_1 + X_2 + \cdots + X_n$. The *sample mean* is

$$M_n = \frac{1}{n} S_n$$

- How good is M_n as an estimator of μ?
  - Use unbiased and consistent to evaluate M_n.

Example: Previous voltage measurement, e.g.,

$$X_i = \mu + N_i$$

where μ is the true value and N_i is the noise.

If we assume N_i are i.i.d. with zero mean,

$$E[X_i] = E[\mu + N_i] = E[\mu] + E[N_i] = \mu + 0 = \mu, \quad \forall i$$

4.3.1   Digression for Sum of RVs (not necessarily independent or identically distributed)

Derivation:

$$E[S_n] = E[X_1 + \cdots + X_n]$$
$$= E[X_1] + \cdots + E[X_n]$$

Derivation:

$$\mathrm{Var}[S_n] = E\left[(S_n - E[S_n])^2\right]$$
$$= E\left[\left(\sum_{i=1}^{n} X_i - E[X_i]\right)^2\right]$$
$$= E\left[\sum_{i=1}^{n}\sum_{j=1}^{n}(X_i - E[X_i])(X_j - E[X_j])\right]$$
$$= \sum_{i=1}^{n}\sum_{j=1}^{n} E\left[(X_i - E[X_i])(X_j - E[X_j])\right]$$
$$= \sum_{i=1}^{n}\sum_{j=1}^{n} \mathrm{Cov}[X_i, X_j]$$
$$= \sum_{i=1}^{n} \mathrm{Var}[X_i] + \sum\sum_{i \neq j} \mathrm{Cov}[X_i, X_j]$$

4.3.2   Unbiased (i.i.d.)

Derivation:

$$E[M_n] = E\left[\frac{1}{n}S_n\right]$$
$$= \frac{1}{n}\left(E[X_1] + \cdots + E[X_n]\right)$$
$$= \frac{1}{n}(n\mu) \quad \text{since } X_i \text{ are i.i.d. so same expectation}$$
$$= \mu \Rightarrow \text{Unbiased!}$$

4.3.3 Consistent (i.i.d.)

Derivation:

$$\text{Var}[M_n] = \text{Var}\left[\frac{1}{n}S_n\right]$$
$$= \frac{1}{n^2}\text{Var}[S_n] \quad \text{taking out constant requires squaring}$$
$$= \frac{1}{n^2}\left(\sum_{i=1}^{n}\text{Var}[X_i] + \sum\sum_{i\neq j}\text{Cov}[X_i, X_j]\right)$$
$$= \frac{1}{n^2}(n\sigma^2) \quad \sigma^2 \triangleq \text{Var}[X_i] \text{ and } X_i \text{ are i.i.d. so covariance is } 0$$
$$= \frac{\sigma^2}{n} \to 0 \text{ as } n \to \infty.$$

- This means that there is no variance in the sample mean as n approaches infinity, so it converges to the true mean.

Recall the Chebyshev Inequality:

$$P\left[|X - E[X]| \geq \epsilon\right] \leq \frac{\text{Var}[X]}{\epsilon^2}, \quad \forall \epsilon > 0.$$

Substitute in M_n:

$$P\left[|M_n - E[M_n]| \geq \epsilon\right] \leq \frac{\text{Var}[M_n]}{\epsilon^2}$$
$$P\left[|M_n - \mu| \geq \epsilon\right] \leq \frac{\sigma^2}{n\epsilon^2}$$
$$\Rightarrow P\left[|M_n - \mu| < \epsilon\right] \geq 1 - \frac{\sigma^2}{n\epsilon^2} \to 1 \text{ as } n \to \infty \text{ then it is consistent}$$

Warning: Cov $= 0$ because independence implies uncorrelated.

4.3.4 Weak Law of Large Numbers

Definition: Even if σ is infinite, then $\forall \epsilon > 0$,

$$\lim_{n\to\infty} P\left[|M_n - \mu| < \epsilon\right] = 1$$

4.3.5   Confidence Interval: Finding n

Example: Measure an unknown voltage θ for n times and obtain independent measurements:

$$X_i = \theta + N_i,$$

where N_i are i.i.d. random variables with mean 0 and variance 1.
- We want to determine how many measurements n are sufficient so that

$$P\left(|M_n - \theta| < 0.1\right) \geq 0.95,$$

  where 0.1 is the desired precision and 0.95 is the confidence level.
- The sample mean is given by:

$$M_n = \frac{1}{n}\sum_{i=1}^{n} X_i = \theta + \frac{1}{n}\sum_{i=1}^{n} N_i.$$

- The variance of X_i is:

$$\sigma^2 = \text{Var}[X_i] = \text{Var}[\theta + N_i] = \text{Var}[N_i] = 1.$$

  - $\text{Var}[aX + b] = a^2\text{Var}[X]$, where $a = 1$ and $b = \theta$.
- Using Chebyshev's inequality:

$$1 - \frac{\sigma^2}{n\epsilon^2} \geq 0.95,$$

  where $\epsilon = 0.1$ (precision).
- Solving for n:

$$1 - \frac{1}{n(0.1)^2} \geq 0.95,$$

$$\frac{1}{n(0.1)^2} \leq 0.05,$$

$$n \geq 2000.$$

Thus, at least 2000 measurements are needed to achieve the desired precision and confidence level.

5   L5: Sample Mean and Maximum Likelihood Estimation

FAQ:
- Why can we say that it is consistent for the last example?

5.1   Maximum Likelihood Estimation

Motivation: Choose parameter θ that is most likely to generate the observation $x_1, x_2, \ldots, x_n$.
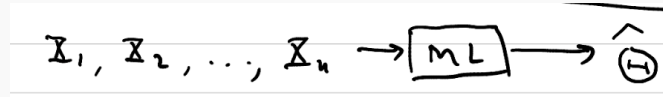
Figure 12:

Definition:

$$\hat{\Theta} = \arg\max_{\theta} P_{\underline{X}}(\underline{x}|\theta), \text{ discrete } X. \tag{41}$$

$$\hat{\Theta} = \arg\max_{\theta} f_{\underline{X}}(\underline{x}|\theta), \text{ continuous } X. \tag{42}$$

5.1.1   Log-Likelihood

Definition:

$$\hat{\theta} = \arg\max_{\theta} \sum_{i=1}^{n} \log P_X(x_i|\theta) \tag{43}$$

$$\hat{\theta} = \arg\max_{\theta} \sum_{i=1}^{n} \log f_X(x_i|\theta). \tag{44}$$

Warning: Can only go from argmax of fcn to argmax of log fcn, if it is i.i.d.

Derivation:

$$\text{i.i.d. } X_1, X_2, \ldots, X_n \implies$$

$$p_{\underline{X}}(\underline{x}|\theta) = \prod_{i=1}^{n} p_{X_i}(x_i|\theta)$$

$$= \prod_{i=1}^{n} p_X(x_i|\theta) \quad \text{drop the i due to i.i.d. assumption}$$

$$\log p_{\underline{X}}(\underline{x}|\theta) = \sum_{i=1}^{n} \log p_X(x_i|\theta).$$

Example:
1. **Model and Observations:**
   - Assume a biased coin with probability θ of showing heads. Find ML estimator for θ.
   - Toss the coin n times and obtain Bernoulli random variables $X_1, \ldots, X_n$ such that:

$$\text{"heads"} \rightarrow 1, \quad \text{"tails"} \rightarrow 0.$$

- Total number of heads is:

$$k = \sum_{i=1}^{n} X_i.$$

  For example:

$$\underline{x} = (1, 0, 0, 0, 1, 1, 0, 1, 1, 1), \quad k = 6.$$

- Probability of observations $x_1, \ldots, x_n$ corresponding to parameter θ is:

$$p_{\underline{X}}(\underline{x}|\theta) = \theta^k (1 - \theta)^{n-k}.$$

  - It is sufficient to know only k.
  - Note: Don't need the $\binom{n}{k}$ term because we are given the specific sequence of heads and tails.

2. **Log-Likelihood and Maximization:**
   - The log-likelihood function is:

$$\log p_{\underline{X}}(\underline{x}|\theta) = k \log(\theta) + (n - k) \log(1 - \theta).$$

   - To maximize the log-likelihood over θ, set:

$$0 = \frac{\partial}{\partial \theta} \log p_{\underline{X}}(\underline{x}|\theta),$$

$$0 = \frac{k}{\theta} - \frac{n - k}{1 - \theta},$$

$$\theta = \frac{k}{n}.$$

   - Thus, the Maximum Likelihood Estimator (MLE) is:

$$\hat{\Theta} = \frac{k}{n}, \quad \text{where } k = \sum_{i=1}^{n} X_i.$$

     This corresponds to the observed frequency of heads, which is intuitive b/c the more heads we see, the more likely the coin is biased towards heads.

3. **Examples:**
   - For $\underline{x} = (1, 0, 0, 0, 1, 1, 0, 1, 1, 1)$:

$$p_X(\underline{x}|\theta) = \theta^6 (1 - \theta)^4,$$

$$\hat{\theta} = \frac{6}{10} = 0.6.$$

   - For $\underline{x} = (0, 1, 1, 1, 0, 0, 1, 0, 1, 0)$:

$$p_X(\underline{x}|\theta) = \theta^5 (1 - \theta)^5,$$

$$\hat{\theta} = \frac{5}{10} = 0.5.$$

Notes:
1. k is a sufficient statistic for this Maximum Likelihood (ML) estimator.

2. The expectation of the estimator $\hat{\theta}$ is:

$$
\begin{aligned}
E[\hat{\Theta}] &= E\left[\frac{1}{n}\sum_{i=1}^{n}X_i\right] \\
&= \frac{1}{n}\sum_{i=1}^{n}E[X_i] \\
&= \frac{1}{n}(n\theta) \\
&= \theta \quad \text{(Unbiased)}.
\end{aligned}
$$

   − $E[X_i] = (1)\theta + (0)(1-\theta) = \theta$

3. In fact, $\hat{\Theta} = \dfrac{1}{n}\sum_{i=1}^{n}X_i$ is the sample mean, and θ is the true mean. Therefore, $\hat{\theta} \to \theta$ in probability, which implies that $\hat{\theta}$ is *consistent*.

6 L6: Maximum Likelihood and Laplace

FAQ:

6.1 MLE for Categorical Random Variables

Example:

1. We say that $X \sim \mathrm{Cat}(\underline{\theta})$ if
$$P[X = m] = \theta_m, \quad m = 1, 2, \ldots, M.$$

   • Going from 2 to M categories is a generalization of the Bernoulli distribution.
   The parameter $\underline{\theta}$ is a vector:
$$\underline{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_M \end{bmatrix},$$

   such that $\theta_m \geq 0$ and $\displaystyle\sum_{m=1}^{M} \theta_m = 1$.

2. Given n i.i.d. observations $X_1, \ldots, X_n$, we aim to find the maximum likelihood estimator (MLE) of $\underline{\theta}$.

3. Define n_m as the number of observations that equal m:
$$n_m = \sum_{i=1}^{n} 1(x_i = m),$$

   where $1(x_i = m)$ is the indicator function. Note that $\displaystyle\sum_{m=1}^{M} n_m = n$.

4. The likelihood function is:
$$p_{\underline{X}}(\underline{x} \mid \underline{\theta}) = \prod_{m=1}^{M} \theta_m^{n_m}.$$

   • Similar to the Bernoulli distribution, but with M categories.
   Taking the log, we get:
$$\log p_{\underline{X}}(\underline{x} \mid \underline{\theta}) = \sum_{m=1}^{M} n_m \log \theta_m.$$

5. To find the optimal $\underline{\theta}$, we minimize the negative log-likelihood:
$$\min_{\underline{\theta}} - \sum_{m=1}^{M} n_m \log \theta_m,$$

   subject to the constraints $\theta_m \geq 0$ for $1 \leq m \leq M$ and $\displaystyle\sum_{m=1}^{M} \theta_m = 1$.

6. Solving this optimization problem, the MLE is:
$$\hat{\Theta}_m = \frac{N_m}{n} = \frac{\sum_{i=1}^{n} 1(X_i = n)}{n}, \quad \underline{\hat{\Theta}} = \begin{bmatrix} \dfrac{N_1}{n} \\ \vdots \\ \dfrac{N_m}{n} \end{bmatrix}.$$

6.2   MLE for Gaussian Random Variables

Example:

1. Given n i.i.d. observations $X_1, \ldots, X_n$ of a Gaussian random variable with parameters (μ, σ^2), we aim to find the maximum likelihood estimators (MLEs) of μ and σ^2.

$$f_{\underline{X}}(\underline{x}|\mu, \sigma^2) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right),$$

$$\log f_{\underline{X}}(\underline{x}|\mu, \sigma^2) = \sum_{i=1}^{n} \left(-\frac{(x_i - \mu)^2}{2\sigma^2} - \log \sigma - \log \sqrt{2\pi}\right).$$

2. To find μ, take the derivative of the log-likelihood with respect to μ and set it to zero:

$$0 = \frac{\partial}{\partial \mu} \sum_{i=1}^{n} \left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right),$$

$$0 = \frac{1}{n} \sum_{i=1}^{n} x_i - \mu,$$

$$\mu = \frac{1}{n} \sum_{i=1}^{n} x_i.$$

3. To find σ^2, take the derivative of the log-likelihood with respect to σ^2 and set it to zero:

$$0 = \frac{\partial}{\partial \sigma^2} \sum_{i=1}^{n} \left(-\frac{(x_i - \mu)^2}{2\sigma^2} - \frac{1}{2}\log \sigma^2\right),$$

$$0 = -\frac{1}{2} \sum_{i=1}^{n} \left(\frac{(x_i - \mu)^2}{\sigma^4}\right) + \frac{1}{2\sigma^2},$$

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^{n} (x_i - \mu)^2.$$

4. Thus, the MLEs are:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^{n} X_i, \quad \text{(sample mean)}$$

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^{n} (X_i - \hat{\mu})^2. \quad \text{(sample variance)}$$

- Note: The sample variance is biased, so we often use $\dfrac{1}{n-1}$ instead of $\dfrac{1}{n}$ to make it unbiased.
- Note: The sample mean is unbiased.

6.3   Will the Sun Rise Tomorrow? (Laplace's Problem)

Example:

- Observation: The Sun has risen for n consecutive days. Estimate the probability that it will rise tomorrow.
- Model: Assume n i.i.d. Bernoulli random variables $X_1, \ldots, X_n$ with $P[X_i = 1] = \theta$.

6.3.1   Frequentist Approach

Example:
1. The Maximum Likelihood Estimator (MLE) is:

$$\hat{\theta} = \frac{K}{n} = \frac{\sum_{i=1}^{n} X_i}{n}.$$

2. If $K = n$ (i.e., the Sun has risen every day so far), then:

$$\hat{\theta} = \frac{n}{n} = 1.$$

3. Conclusion: The Sun will rise tomorrow with probability 1, regardless of what n is, based on the Frequentist approach.
   - This doesn't make sense b/c if $n = 1$ then we are assuming 100% it will rise based on one observation.

6.3.2   Bayesian Approach

Example:
1. Assume that θ is not fixed but drawn from a uniform distribution in $[0, 1]$. This means that the probability of the Sun rising is based on a uniform distribution.
2. We want to find the probability that the sun will rise tomorrow given that it has risen for n consecutive days:

$$P[X_{n+1} = 1 | X_1 = 1, \ldots, X_n = 1].$$

Using Bayes' Theorem:

$$P[X_{n+1} = 1 | X_1 = 1, \ldots, X_n = 1] = \frac{P[X_1 = 1, \ldots, X_{n+1} = 1]}{P[X_1 = 1, \ldots, X_n = 1]}.$$

3. Compute $P[X_1 = 1, \ldots, X_n = 1]$:

$$P[X_1 = 1, \ldots, X_n = 1] = \int_0^1 P[X_1 = 1, \ldots, X_n = 1 | \Theta = \theta] f_\Theta(\theta) \, d\theta.$$

   - The joint probability is calculated by integrating the product of the likelihood and the prior (i.e. marginalizing over θ).
   - The likilihood becomes θ^n because the observations are i.i.d.
   - The prior is uniform, so $f_\Theta(\theta) = 1$.

   Since $f_\Theta(\theta) = 1$ (uniform prior) and $P[X_1 = 1, \ldots, X_n = 1 | \Theta = \theta] = \theta^n$, we have:

$$P[X_1 = 1, \ldots, X_n = 1] = \int_0^1 \theta^n \, d\theta = \frac{1}{n+1}.$$

4. Compute $P[X_1 = 1, \ldots, X_{n+1} = 1]$ similarly:

$$P[X_1 = 1, \ldots, X_{n+1} = 1] = \int_0^1 \theta^{n+1} \, d\theta = \frac{1}{n+2}.$$

5. Combine results:

$$P[X_{n+1} = 1 | X_1 = 1, \ldots, X_n = 1] = \frac{\frac{1}{n+2}}{\frac{1}{n+1}} = \frac{n+1}{n+2}.$$

6. Conclusion: As n increases, the probability approaches 1, providing more certainty with more data.

6.4   Sample Mean is Not Always an ML Estimator

Example: Given an unknown voltage x, we measure it using a voltmeter that outputs a random reading Y that is uniform in $[0, x]$. Suppose we make n i.i.d. measurements $Y_1, \ldots, Y_n$ and wish to estimate $\mu = \dfrac{x}{2} = \mathbb{E}[Y]$.

1. PDF of Y:

$$f_Y(y \mid \mu) = \begin{cases} \dfrac{1}{2\mu} & \text{if } 0 \leq y \leq 2\mu, \\ 0 & \text{otherwise.} \end{cases}$$

   - This is a uniform distribution, where $x = 2\mu$.

2. **Sample Mean:**

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^{n} Y_i \quad \cdots (1)$$

3. **To Find the ML Estimator:**

$$f_{\mathbf{Y}}(\mathbf{y} \mid \mu) = \prod_{i=1}^{n} f_Y(y_i \mid \mu)$$

$$= \prod_{i=1}^{n} \frac{1}{2\mu} \cdot 1(0 \leq y_i \leq 2\mu)$$

$$= \frac{1}{(2\mu)^n} \prod_{i=1}^{n} 1(0 \leq y_i \leq 2\mu).$$

   - The indicator function ensures that all measurements are within the range $[0, 2\mu]$ for the likelihood to be non-zero. This is done because we assume that Y is uniformly distributed in $[0, 2\mu]$.

   The likelihood is maximized for:

$$\arg\max_{\mu} f_{\mathbf{Y}}(\mathbf{y} \mid \mu) = \max_{1 \leq i \leq n} \frac{1}{2} Y_i.$$

   Therefore:

$$\hat{\mu} = \max_{1 \leq i \leq n} \frac{1}{2} Y_i \quad \cdots (2)$$

   - The likelihood is non-zero only if μ is greater than or equal to the maximum of the measurements because all data points must lie within the range $[0, 2\mu]$.
   - i.e. $2\mu \geq \max\limits_{1 \leq i \leq n} Y_i$, which is to ensure that ALL measurements are within the range $[0, 2\mu]$, so this must be μ.

4. Clearly, $(1) \neq (2)$.

7   L7: Maximum A Posteriori (MAP) Estimation

FAQ:

7.1   ML Estimator (Frequentist Approach)

Definition:
- Assume θ is **fixed** and find the "best" θ:

$$\hat{\theta} = \arg\max_{\theta} P_{\underline{X}}(\underline{x}|\theta) \quad \text{or} \quad \hat{\theta} = \arg\max_{\theta} f_{\underline{X}}(\underline{x}|\theta).$$

- Every θ value specifies a **different probability space** (e.g., coin, universe, etc.) for our experiment.

7.2   MAP Estimator (Bayesian Approach)

Definition:
1. Assume **random** parameter Θ in the **same sample space** as our experiment.
2. Assume we have a **prior** pmf/pdf for Θ: $P_{\Theta}(\theta)$ or $f_{\Theta}(\theta)$.
3. Find the most probable θ given the observations $\underline{X} = \underline{x}$:

$$\hat{\theta} = \arg\max_{\theta} P_{\Theta|\underline{X}}(\theta|\underline{x}), \quad \text{if } \theta \text{ discrete},$$

$$\hat{\theta} = \arg\max_{\theta} f_{\Theta|\underline{X}}(\theta|\underline{x}), \quad \text{if } \theta \text{ continuous}.$$

- **Posterior Distribution:** $f_{\Theta|\underline{X}}(\theta|\underline{x})$.

7.2.1   Four Cases of Bayes' Rule

Definition:

$$P_{\Theta|\underline{X}}(\theta|\underline{x}) = \begin{cases} \dfrac{P_{\underline{X}|\Theta}(\underline{x}|\theta)P_{\Theta}(\theta)}{P_{\underline{X}}(\underline{x})}, & \text{if } \underline{X} \text{ discrete}, \\ \dfrac{f_{\underline{X}|\Theta}(\underline{x}|\theta)P_{\Theta}(\theta)}{f_{\underline{X}}(\underline{x})}, & \text{if } \underline{X} \text{ continuous}. \end{cases}$$

$$f_{\Theta|\underline{X}}(\theta|\underline{x}) = \begin{cases} \dfrac{P_{\underline{X}|\Theta}(\underline{x}|\theta)f_{\Theta}(\theta)}{P_{\underline{X}}(\underline{x})}, & \text{if } \underline{X} \text{ discrete}, \\ \dfrac{f_{\underline{X}|\Theta}(\underline{x}|\theta)f_{\Theta}(\theta)}{f_{\underline{X}}(\underline{x})}, & \text{if } \underline{X} \text{ continuous}. \end{cases}$$

Notes: The denominator $f_{\underline{X}}(\underline{x})$ is independent of θ:

$$f_{\underline{X}}(\underline{x}) = \int_{-\infty}^{\infty} f_{\underline{X}|\Theta}(\underline{x}|\theta)f_{\Theta}(\theta)d\theta.$$

7.3   Comparison: ML vs MAP

Definition:

$$\hat{\theta}_{ML} = \arg\max_{\theta} P_{\underline{X}}(\underline{x}|\theta),$$

$$\hat{\theta}_{MAP} = \arg\max_{\theta} P_{\underline{X}|\Theta}(\underline{x}|\theta)P_{\Theta}(\theta).$$

The expressions are algebraically similar, but the philosophies differ.

7.4   Example: Unknown Voltage (REVIEW)

Example:
1. For an unknown voltage, we have some prior knowledge that θ is uniformly distributed in $[0, 1]$:

$$f_\Theta(\theta) = \begin{cases} 1, & 0 \le \theta \le 1, \\ 0, & \text{otherwise.} \end{cases}$$

2. We measure it using a voltmeter that outputs random readings Y that is uniformly distributed in $[0, \theta]$. Suppose we make n measurements $\underline{Y} = (Y_1, Y_2, \ldots, Y_n)$.
3. How to estimate θ?

To Warm Up: ML Estimation (Ignoring the Prior)

$$f_{\underline{Y}|\Theta}(\underline{y}|\theta) = \prod_{i=1}^{n} f_{Y|\Theta}(y_i|\theta)$$

$$= \frac{1}{\theta^n} \prod_{i=1}^{n} 1(0 \le y_i \le \theta)$$

$$= \frac{1}{\theta^n} 1(\theta \ge \max_{1 \le i \le n} y_i) \quad (\text{since } 0 \le y_i \le \theta).$$

$$\hat{\theta}_{ML} = \arg\max_\theta \frac{1}{\theta^n} 1(\theta \ge \max_{1 \le i \le n} y_i) = \max_{1 \le i \le n} y_i.$$

- Since we want smallest θ to maximize the likelihood, we choose the largest y_i (i.e. lowest bound)

Solution: MAP Estimation

$$f_{\Theta|\underline{Y}}(\theta|\underline{y}) = \frac{f_{\underline{Y}|\Theta}(\underline{y}|\theta) f_\Theta(\theta)}{f_{\underline{Y}}(\underline{y})}$$

$$\propto f_{\underline{Y}|\Theta}(\underline{y}|\theta) f_\Theta(\theta)$$

$$= \frac{1}{\theta^n} 1(\theta \ge \max_{1 \le i \le n} y_i) 1(0 \le \theta \le 1).$$

$$\hat{\theta}_{MAP} = \arg\max_\theta \frac{1}{\theta^n} 1(\theta \ge \max_{1 \le i \le n} y_i) 1(0 \le \theta \le 1) = \max_{1 \le i \le n} y_i.$$

- Since we know $0 \le y_i \le 1$, therefore, choose the largest y_i (i.e. lowest bound)

What If the Prior Is Different? Suppose the prior is:

$$f_\Theta(\theta) = \begin{cases} 2, & \frac{1}{2} \le \theta \le 1, \\ 0, & \text{otherwise.} \end{cases}$$

This means "The voltage is at least $\frac{1}{2}$ but no more than 1." Then:

$$f_{\Theta|\underline{Y}}(\theta|\underline{y}) \propto \frac{1}{\theta^n} 1(\theta \ge \max_{1 \le i \le n} y_i) 2 \cdot 1\left(\frac{1}{2} \le \theta \le 1\right),$$

$$\hat{\theta}_{MAP} = \max\left\{\max_{1 \le i \le n} y_i, \frac{1}{2}\right\}.$$

Notes:
1. $\max_{1 \le i \le n} y_i$ is a sufficient statistic for ML and MAP in this scenario.
2. $\hat{\theta}_{MAP} \to \max_{1 \le i \le n} y_i$ as $n \to \infty$, regardless of the prior.
3. MAP optimization typically requires more computational effort than ML estimation.

8 L8: MAP Conjugate Prior

8.1 Example: Coin Tosses

Example:
- Let $\theta = P(\text{Heads})$.
- Let X be the number of heads from n independent tosses.
- Given observation X, find the MAP (Maximum a Posteriori) estimate for θ.

The likelihood is given by:

$$P_{X|\theta}(k|\Theta) = \binom{n}{k}\theta^k(1-\theta)^{n-k}.$$

The posterior distribution is proportional to:

$$f_{\Theta|X}(\theta|k) \propto \binom{n}{k}\theta^k(1-\theta)^{n-k}f_\Theta(\theta).$$

8.1.1 MAP Estimation

Example: We will consider only integer α and β.

$$f_{\Theta|X}(\theta|k) = \frac{\binom{n}{k}\theta^k(1-\theta)^{n-k} \cdot B(\alpha,\beta)\theta^{\alpha-1}(1-\theta)^{\beta-1}}{f_X(k)}$$

$$\propto \theta^{k+\alpha-1}(1-\theta)^{n-k+\beta-1}.$$

To normalize the posterior:

$$f_{\theta|X}(\theta|k) = \frac{1}{B(k+\alpha, n-k+\beta)}\theta^{k+\alpha-1}(1-\theta)^{n-k+\beta-1}.$$

- i.e. To have the integrating PDF equal to 1, we need to divide by $B(k+\alpha, n-k+\beta)$.

The MAP (Maximum a Posteriori) estimate is given by:

$$\hat{\theta}_{\text{MAP}} = \arg\max_\theta f_{\Theta|X}(\theta|k)$$

$$= \frac{k+\alpha-1}{n+\alpha+\beta-2}.$$

Notes:
1. Without using any prior, the MLE (Maximum Likelihood Estimate) is:

$$\hat{\theta}_{ML} = \frac{k}{n}.$$

2. For $\alpha = \beta = 1$ (uniform prior), $\hat{\theta}_{MAP} = \hat{\theta}_{ML}$.
3. For general α and β, we can interpret the β prior as conducting $\alpha + \beta - 2$ prior coin tosses and observing $\alpha - 1$ heads.
4. As $n \to \infty$, the impact of the prior diminishes, and $\hat{\theta}_{MAP} \to \hat{\theta}_{ML}$.

8.1.2 What is a good choice for $f_\Theta(\theta)$?

Motivation: We aim for a prior that provides a rich representation of prior belief/information and facilitates numerical calculation and optimization.

8.2   Beta Prior

Definition: Consider the **Beta prior**, where Θ is a Beta random variable with parameters $\alpha > 0$ and $\beta > 0$:

$$f_\Theta(\theta) = \begin{cases} \dfrac{1}{B(\alpha, \beta)} \theta^{\alpha-1}(1-\theta)^{\beta-1}, & 0 < \theta < 1, \\ 0, & \text{otherwise.} \end{cases}$$

Here, $B(\alpha, \beta)$ is the Beta function.

8.2.1   Digression: Beta Function and Beta Random Variables

Definition:
- The Beta function can also be expressed in terms of the Gamma function:

$$B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)},$$

where $\Gamma(x) = \displaystyle\int_0^\infty t^{x-1} e^{-t}\, dt$ is the Gamma function.

Derivation:

$$1 = \int_0^1 f_\Theta(\theta)\, d\theta = \int_0^1 \frac{1}{B(\alpha, \beta)} \theta^{\alpha-1}(1-\theta)^{\beta-1}\, d\theta,$$

$$\to B(\alpha, \beta) = \int_0^1 \theta^{\alpha-1}(1-\theta)^{\beta-1}\, d\theta.$$

8.2.2   Properties of the Beta Distribution

Definition: Properties of the Gamma function:
1. $\Gamma(x + 1) = x\Gamma(x)$. For integer m, $\Gamma(m + 1) = m!$.
2. For integers α and β:
$$B(\alpha, \beta) = \frac{(\alpha + \beta - 1)!}{(\alpha - 1)!(\beta - 1)!} = \beta\binom{\alpha + \beta - 1}{\alpha - 1}.$$
3. Expected Value of a Beta Random Variable
$$E[\theta] = \int_0^1 \theta \frac{1}{B(\alpha, \beta)} \theta^{\alpha-1}(1-\theta)^{\beta-1}\, d\theta = \frac{\alpha}{\alpha + \beta}.$$
4. The max of Beta PDF (i.e. mode) is $\theta = \dfrac{\alpha - 1}{\alpha + \beta - 2}$.

Derivation:

1. Expectation:

$$
\begin{aligned}
E[\Theta] &= \int_0^1 \theta \cdot \frac{1}{B(\alpha,\beta)} \theta^{\alpha-1}(1-\theta)^{\beta-1}\, d\theta \\
&= \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \int_0^1 \theta^{\alpha}(1-\theta)^{\beta-1}\, d\theta \\
&= \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} B(\alpha+1,\beta) \\
&= \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \cdot \frac{\Gamma(\alpha+1)\Gamma(\beta)}{\Gamma(\alpha+\beta+1)} \\
&= \frac{\Gamma(\alpha+1)\Gamma(\beta)}{\Gamma(\alpha)\Gamma(\beta)} \cdot \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha+\beta+1)} \\
&= \frac{\alpha}{\alpha+\beta}.
\end{aligned}
$$

2. Mode: To find the mode, take the derivative with respect to θ:

$$
\log f_\theta(\theta) = \log \frac{1}{B(\alpha,\beta)} + (\alpha-1)\log\theta + (\beta-1)\log(1-\theta).
$$

$$
\begin{aligned}
0 &= \frac{d}{d\theta}\log f_\theta(\theta) \\
&= (\alpha-1)\frac{1}{\theta} + (\beta-1)\cdot\frac{-1}{1-\theta}.
\end{aligned}
$$

Simplify and solve for θ:

$$
\theta = \frac{\alpha-1}{\alpha+\beta-2}, \quad \text{for } \alpha,\beta > 1.
$$

8.2.3   Visualization of the Beta Prior

Notes:
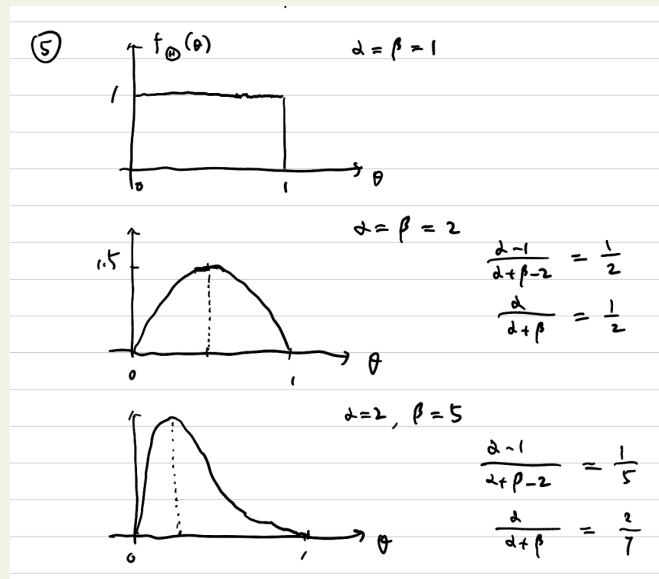- Larger α and β imply greater certainty in the prior.

Figure 13:

8.3   Conjugate Priors

Notes: For ease of algebraic manipulation, we prefer priors $f_\Theta(\theta)$ with the same functional form as the posterior $f_{\Theta|\underline{X}}(\theta|\underline{x})$. Examples:
- Beta prior for Bernoulli, binomial, and geometric distributions.
- Dirichlet prior for categorical and multinomial distributions.
- Gamma prior for Poisson, exponential, and Gaussian distributions.
- Gaussian prior for Gaussian distributions.

9   L9: Least Mean Squares (LMS) Estimation

Definition: **Assume prior:** $P_\Theta(\theta)$ or $f_\Theta(\theta)$ w/ observations: $\bar{X} = x$.

$$\hat{\theta} = g(\underline{x}) = \mathbb{E}[\Theta \mid \underline{X} = \underline{x}]$$
$$\text{or } \hat{\Theta} = g(\underline{X}) = \mathbb{E}[\Theta \mid \underline{X}].$$

Notes:
1. **MAP vs. LMS Estimators:**
    - **MAP:** Use the most probable θ given x.
    - **LMS:** Use the expected value (conditional on $\bar{X} = x$) of Θ, i.e., the "Conditional Expectation Estimator."
2. **Unbiasedness of LMS Estimator:**

$$\mathbb{E}[\hat{\Theta}] = \mathbb{E}[\mathbb{E}[\Theta \mid \underline{X}]] = \mathbb{E}[\Theta],$$
$$\implies \mathbb{E}[\hat{\Theta} - \Theta] = 0.$$

3. **LMS Estimator Minimizes Conditional MSE:**

$$\mathbb{E}\big[(\Theta - \hat{\Theta})^2 \mid \underline{X} = \underline{x}\big].$$

Proof:
(a) First, suppose no observations: $\hat{\Theta}$ is a constant. So we want:

$$\hat{\Theta} = \arg\min_c \mathbb{E}\big[(\Theta - c)^2\big],$$
$$0 = \frac{d}{dc}\big[-2\mathbb{E}[\Theta] + 2c\big],$$
$$c = \mathbb{E}[\Theta].$$

(b) Alternate view:

$$\mathbb{E}\big[(\Theta - c)^2\big] = \text{Var}[\Theta - c] + \mathbb{E}[\Theta - c]^2,$$
$$= \text{Var}[\Theta] + (\mathbb{E}[\Theta] - c)^2.$$

To minimize: Set bias $\mathbb{E}[\Theta] - c$ to zero, while for variance, we have no control.
(c) Now, with observations $\underline{X} = \underline{x}$ (i.e. given): $\hat{\theta} = g(\underline{x})$

$$\mathbb{E}\big[(\Theta - g(\underline{x}))^2 \mid \underline{X} = \underline{x}\big] = \text{Var}[\Theta \mid \underline{X} = \underline{x}] + (\mathbb{E}[\Theta \mid \bar{X} = \underline{x}] - g(\underline{x}))^2.$$

To minimize: Set $(g(x) - \mathbb{E}[\Theta \mid \underline{X} = \underline{x}])^2 = 0$ since we have no control over the variance.
(d) Conclusion:
$$\hat{\theta} = g(x) = \mathbb{E}[\Theta \mid \underline{X} = \underline{x}].$$

9.1   Example: Prior Coin Toss Problem

Example:

$$\hat{\theta}_{\text{LMS}} = \mathbb{E}[\Theta \mid X = k]$$
$$= \frac{k + \alpha}{n + \alpha + \beta}.$$

9.2　Example: Prior Voltage Problem

Example:
1. **Setup:**
    - Unknown voltage Θ.
    - Prior: $\Theta \sim \text{Uniform}[0, 1]$.
    - Volt meter reading Y given $\Theta = \theta$: $Y \sim \text{Uniform}[0, \Theta]$.
    - Independent measurements: $Y_1, \ldots, Y_n$ given θ.
2. **Likelihood:**

$$f_{\underline{Y}|\Theta}(\underline{y} \mid \theta) = \prod_{i=1}^{n} f_{Y|\Theta}(y_i \mid \theta)$$
$$= \frac{1}{\theta^n} \cdot 1(\theta \geq \max_{1 \leq i \leq n} y_i).$$

3. **Posterior:**

$$f_{\Theta|\underline{Y}}(\theta \mid \underline{y}) = \frac{\frac{1}{\theta^n} \cdot \mathbf{1}(\theta \geq \max_{1 \leq i \leq n} y_i)\mathbf{1}(0 \leq \theta \leq 1)}{f_{\underline{Y}}(\underline{y})}.$$

4. **Estimators:**
    - ML/MAP:
$$\hat{\theta} = \max_{1 \leq i \leq n} y_i.$$
    - LMS:
$$\hat{\theta} = \mathbb{E}[\Theta \mid \underline{Y} = \underline{y}] = \int_{-\infty}^{\infty} \theta f_{\Theta|\underline{Y}}(\theta \mid \underline{y})d\theta.$$

5. **Derivation for LMS:**
    (a) We need $f_{\underline{Y}}(\underline{y}) = \int_{-\infty}^{\infty} \frac{1}{\theta^n} \cdot \mathbf{1}(\theta \geq \max_{1 \leq i \leq n} y_i)\mathbf{1}(0 \leq \theta \leq 1)d\theta$
    (b) Compute $f_{\underline{Y}}(\underline{y})$ for $n = 1$:

$$f_{\underline{Y}}(\underline{y}) = \int_{y}^{1} \frac{1}{\theta}d\theta, \quad 0 \leq y \leq 1$$
$$= \ln(\theta)\Big|_{y}^{1}$$
$$= -\ln(y).$$

    (c) Compute $\hat{\theta}$ for $n = 1$:

$$\hat{\theta} = \int_{-\infty}^{\infty} \theta \frac{\frac{1}{\theta}\mathbf{1}(\theta \geq y) \cdot \mathbf{1}(0 \leq \theta \leq 1)}{-\ln y}d\theta$$
$$= \frac{1}{-\ln(y)} \int_{y}^{1} d\theta$$
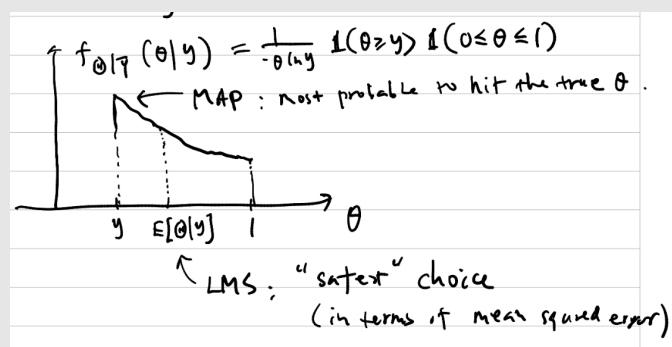$$= \frac{y - 1}{ln(y)}$$

6. **Graphical Interpretation:**

$$f_{\Theta|Y}(\theta|y) = \frac{1}{-\theta \ln y} \mathbb{1}(\theta \geq y) \mathbb{1}(0 \leq \theta \leq 1)$$

MAP : Most probable to hit the true θ.

$y \quad E[\theta|y] \quad 1 \qquad \theta$

LMS : "safer" choice
(in terms of mean squared error)

Figure 14: