ECE353 Lectures

Hanhee Lee

January 15, 2025

Contents

1 Prologue **3**
1.1 Components of a Robotic System . . . 3
1.1.1 Overview (Robots, the Environment) . . . 4
1.1.2 Robot (Sensors, Actuators, the Brain) . . . 4
1.1.3 Brain (Tracker, Planner, Memory) . . . 5
1.1.4 Environment (Physics, State) . . . 5
1.2 Equations of a Robotic System . . . 6
1.2.1 Sensing . . . 6
1.2.2 Tracking . . . 6
1.2.3 Planning . . . 6
1.2.4 Acting . . . 7
1.2.5 Simulating . . . 7
1.3 Setup of Planning Problems . . . 8

2 Uninformed Search Algorithms **8**
2.1 Setup . . . 8
2.2 Search Graphs . . . 8
2.3 Path Trees . . . 8
2.4 Search Algorithms . . . 9
2.4.1 Characteristics of a Search Algorithm . . . 9
2.4.2 Breadth First Search (BFS) . . . 9
2.4.3 Depth First Search (DFS) . . . 10
2.4.4 Iterative Deepening DFS (IDDFS) . . . 10
2.4.5 Cheapest-First Search (CFS) . . . 10
2.5 Modifications to Search Algorithms . . . 10
2.5.1 Depth-Limiting . . . 10
2.5.2 Iterative Deepening . . . 11
2.5.3 Cost-Limiting . . . 11
2.5.4 Iterative-Inflating . . . 11
2.5.5 Intra-Path Cycle Checking . . . 12
2.5.6 Inter-Path Cycle Checking . . . 12
2.6 Informed Search Algorithms . . . 12
2.6.1 Estimated Cost . . . 12
2.7 Characteristics of an Informed Search Algorithm . . . 12
2.7.1 Heuristics . . . 13
2.7.2 Heuristic-First Search (HFS) . . . 13
2.7.3 A-Star Search (A*) . . . 13
2.7.4 Iterative Inflating A-Star Search (IIA*) . . . 13
2.7.5 Designing Heuristics via Problem Relaxation . . . 13
2.7.6 Combining Heuristics . . . 13
2.8 Anytime Search Algorithms . . . 13
2.9 Formulating a Search Problem . . . 13

3 Informed Search Algorithms **13**

4 Probability Review **13**

5 Probabilistic Inference Problems **13**

1 Prologue

Summary:
- This course will focus on planning
- Variables:
  - State: $\mathbf{x}(t)$
  - Action(s): $\mathbf{u}(t)$
  - Measurement: $\mathbf{y}_k^{(i)}$
  - Context: $\mathbf{z}_k^{(i)}$
  - Old Context: $\mathbf{z}_{k-1}^{(i)}$
  - Plan: $\mathbf{p}_k^{(i)}$
  - (i): Ith agent
- Conversion to DT is necessary because robots are digitalized system and then converted back to CT for execution.

FAQ:
- What does the environment do?
- What is the joint action set?

1.1 Components of a Robotic System

Summary:

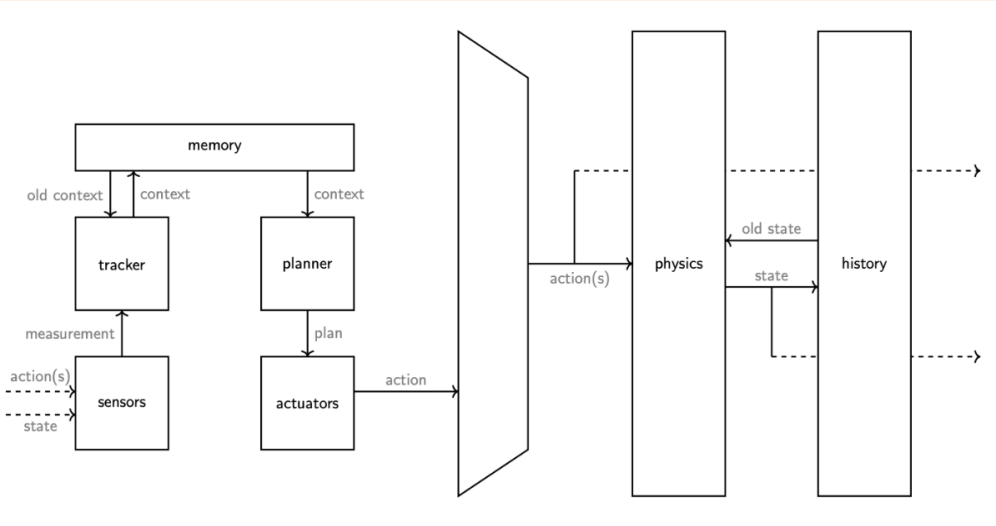

Figure 1: Components of a Robotic System (Words)

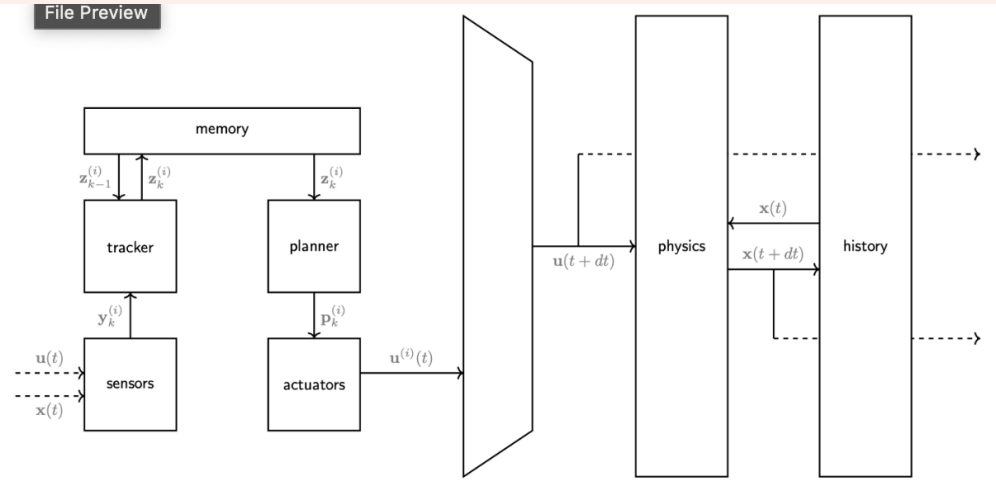

Figure 2: Components of a Robotic System (Math)

1.1.1 Overview (Robots, the Environment)

Definition:

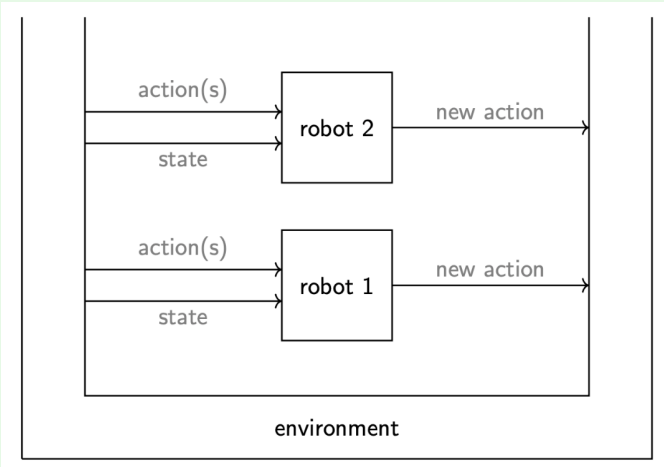

Figure 3: Overview (Robots, the Environment)

Notes:
- Environment → previous actions + current state → robot → new action → environment

1.1.2 Robot (Sensors, Actuators, the Brain)

Definition:

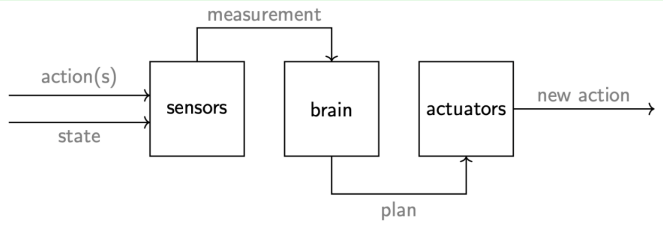

Figure 4: Robot (Sensors, Actuators, the Brain)

Notes:
- Measurements can be noisy and inaccurate if not a perfect sensor.
- Measurements go into the brain which can create a plan.

1.1.3 Brain (Tracker, Planner, Memory)

Definition:

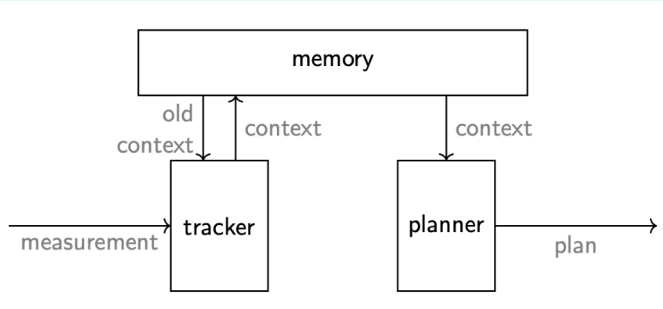

Figure 5: Brain (Tracker, Planner, Memory)

Notes:
- The tracker takes in the measurements and old context and updates the context.
- The planner takes in the context and creates a plan.
- The memory stores the context.

1.1.4 Environment (Physics, State)

Definition:

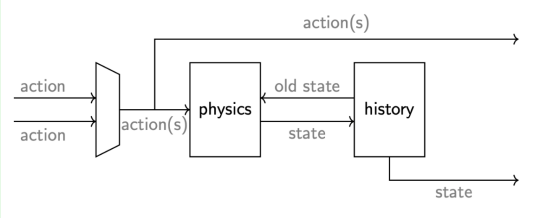

Figure 6: Environment (Physics, State)

1.2 Equations of a Robotic System

1.2.1 Sensing

Definition: Take a measurement:

$$\mathbf{y}^{(i)}(t) = \mathrm{sns}^{(i)}\big(\mathbf{x}(t), \mathbf{u}(t), t\big)$$

Convert the measurement into a discrete-time signal using a sampling period of $T^{(i)}$:

$$\mathbf{y}_k^{(i)} = \mathrm{dt}\big(\mathbf{y}^{(i)}(t), t, T^{(i)}\big)$$

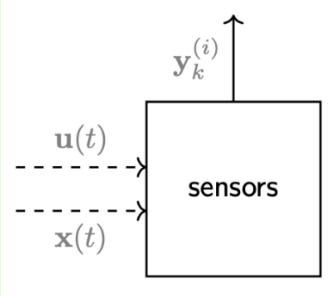

Figure 7: Sensing

1.2.2 Tracking

Definition: Track (update) the context:

$$\mathbf{z}_k^{(i)} = \mathrm{trk}^{(i)}\big(\mathbf{z}_{k-1}^{(i)}, \mathbf{y}_k^{(i)}, k\big)$$

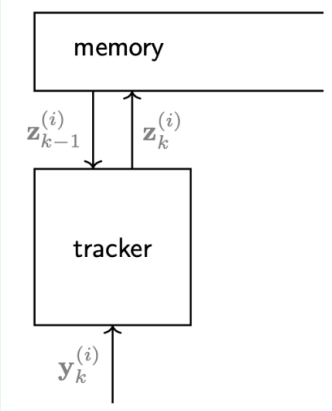

Figure 8: Tracking

1.2.3 Planning

Definition: Make a plan:

$$\mathbf{p}_k^{(i)} = \mathrm{pln}^{(i)}\big(\mathbf{z}_k^{(i)}, k\big)$$

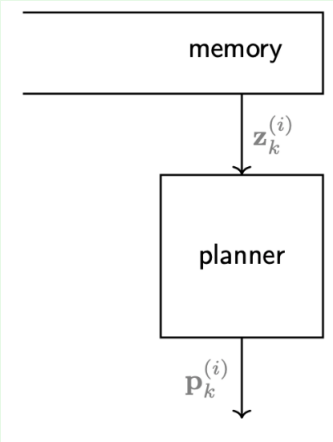

Figure 9: Planning

1.2.4 Acting

Definition: Convert the plan into a continuous-time signal using a sampling period of $T^{(i)}$:

$$\mathbf{p}(t) = \mathrm{ct}\big(\mathbf{p}_k^{(i)}, t, T^{(i)}\big)$$

Execute the plan:

$$\mathbf{u}^{(i)}(t) = \mathrm{act}^{(i)}\big(\mathbf{p}^{(i)}(t), t\big)$$

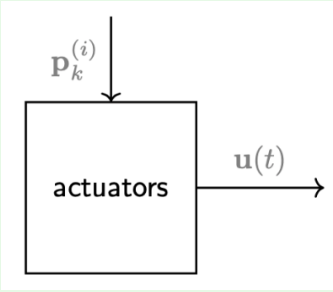

Figure 10: Acting

1.2.5 Simulating

Definition: Simulate the environment's response:

$$\dot{\mathbf{x}}(t) = \mathrm{phy}\big(\mathbf{x}(t), \mathbf{u}(t), t\big)$$

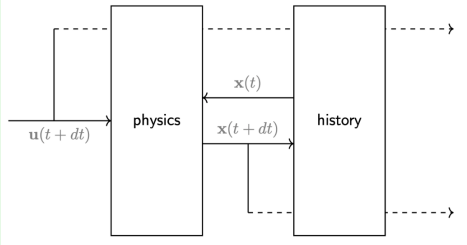

Figure 11: Simulating

1.3 Setup of Planning Problems

Summary: In a planning problem, it is assumed that:
- the environment is representable using a discrete set of states, $\mathcal{S}$
- for each state, $s \in \mathcal{S}$, each agent, i, has a discrete set of actions, $\mathcal{A}_i(s)$, with $\mathcal{A}(s) := \times_i \mathcal{A}_i(s)$
- a **move** is any tuple, (s, a), where $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$
- a **transition** is any 3-tuple, (s, a, s'), where $s, s' \in \mathcal{S}$ and $a \in \mathcal{A}(s)$
- the transition resulting from a move may be deterministic/stochastic
- $rwd_i(s, a, s')$ is agent i's reward for the transition, (s, a, s')
- a **path** is any sequence of transitions of the form

$$p = \left\langle (s^{(0)}, a^{(1)}, s^{(1)}), (s^{(1)}, a^{(2)}, s^{(2)}), \dots \right\rangle$$

- each agent wants to realize a path that maximizes its own reward

Warning: $\mathcal{A}(s)$ is the joint action set of all agents at state s.

2 Uninformed Search Algorithms

Summary:
- Not responsible for proofs, but know when to use each algorithm.

2.1 Setup

Definition: In a search problem, it is assumed that:
- There is only one agent (us).
- For each state, $s \in S$, we have a discrete set of actions, $\mathcal{A}(s)$.
- The transition resulting from a move, (s, a), is deterministic; the resulting state is $tr(s, a)$.
- $cst(s, a, tr(s, a))$ is our cost for the transition, $(s, a, tr(s, a))$.
- We want to realize a path that minimizes our cost.

A search problem may have no solutions, in which case, we define the solution as `NULL`.

2.2 Search Graphs

Definition: In a search graph (a graph representing a search problem):
- S is defined by the vertices.
- $\mathcal{G}$ is a subset of the vertices.
- $s^{(0)}$ is some vertex.
- $tr(\cdot, \cdot)$ and $\mathcal{T}$ are defined by the edges.
- $cst(\cdot, \cdot, \cdot)$ is defined by the edge weights.

2.3 Path Trees

Definition: A search algorithm explores a tree of possible paths.
- In such a tree, each node represents the path from the root to itself.
  - The node may also include other info (such as the path's origiin, cost, etc).

2.4 Search Algorithms

Definition: All search algorithms follow the template below:

```
1 O ← {(⟨⟩, 0)}                                          ▷ initialize a set of open nodes
2 SEARCH(O)
```

- $\langle\rangle$ is the empty path, and 0 is the cost of the empty path.

```
1 procedure SEARCH(O)
2     if O = ∅ then
3         return NULL                                    ▷ the search algorithm failed to find a path to a goal
4     n ← REMOVE(O)                                      ▷ "explore" a node n
5     if DST(n) ∈ G then
6         return n                                       ▷ the search algorithm found a path to a goal
7     for n' ∈ CHL(n) do
8         O ← O ∪ {n'}                                   ▷ "expand" n and "export" its children
9     SEARCH(O)
```

- Explore: Remove a node from the open set.
- Exapnd: Generate the children of the node.
- Export: Add the children to the open set.

Warning: The key difference is in the order that REMOVE$(\cdot)$ removes nodes.

2.4.1 Characteristics of a Search Algorithm

Definition: We want to choose `REMOVE(·)` so that the algorithm exhibits the following characteristics:

Characteristic	Description
Halting	Terminates after finitely many nodes explored
Sound	Returned (possibly NULL) solution is correct
Complete	Halting and sound when a non-NULL solution exists
Optimal	Returns an optimal solution when multiple exist
Time Efficient	Minimizes the nodes **explored**/expanded/exported
Space Efficient	Minimizes the nodes simultaneously open

- Will be using explored for time efficiency.

The characteristics of the algorithm also depend on several properties of the path tree over which it searches. These properties include:

- Branching factor: b ($b < \infty$), the maximum number of children a node can have.
- Depth: d, the length of the longest path.
- Length of the shortest solution: l^*
- Cost of the cheapest solution: c^*
- Cost of the cheapest edge: ϵ

We want to choose `REMOVE(·)` so that the algorithm exhibits the aforementioned characteristics for as many path trees as possible.

2.4.2 Breadth First Search (BFS)

Definition: Explores the least-recently expanded open node first.

Property	Description
Halting	$d < \infty$ non-NULL
Sound	always
Complete	always
Optimal	constant cst
Time	b^{l^*}
Space	b^{l^*+1}

2.4.3 Depth First Search (DFS)

Definition: Explores the most-recently expanded open node first.

Property	Description
Halting	$d < \infty$
Sound	always
Complete	$d < \infty$
Optimal	never
Time	b^d
Space	bd

2.4.4 Iterative Deepening DFS (IDDFS)

Definition: Same as DFS but with iterative deepening.

Property	Description
Halting	always
Sound	always
Complete	always
Optimal	constant cst
Time	b^{l^*}
Space	bl^*

2.4.5 Cheapest-First Search (CFS)

Definition: Explores the cheapest open node first.

Property	Description
Halting	$d < \infty$ non-NULL
Sound	yes
Complete	$\epsilon > 0$
Optimal	$\epsilon > 0$
Time	$b^{c^*/\epsilon}$
Space	$b^{c^*/\epsilon+1}$

2.5 Modifications to Search Algorithms

2.5.1 Depth-Limiting

Definition: Depth limit of $d_{\max}$ to any search algorithm by modifying `SEARCH(·)` as follows:

```
procedure SEARCHDL(𝒪, d_max):
    if 𝒪 = ∅ then
        return NULL                                  ▷ the search algorithm failed to find a path to a goal
    n ← REMOVE(𝒪)                                    ▷ "explore" a node, n
    if dst(n) ∈ 𝒢 then
```

```
        return n                                          ▷ the search algorithm found a path to a goal
    for n' ∈ chl(n) do                                    ▷ "expand" n and "export" its children
        if len(n') ≤ d_max then                           ▷ unless the child is too long
            O ← O ∪ {n'}
    SEARCHDL(O, d_max)
```

2.5.2 Iterative Deepening

Definition: Iteratively increase the depth-limit, $d_{\max}$, to any search algorithm w/ depth-limiting, by placing `SEARCHDL(·)` in a wrapper, `SEARCHID(·)`:

```
procedure SEARCHID():
    n ← NULL
    d_max ← 0
    ▷ while a solution has not been found, reset the open set, run the search algorithm, then increase the
    depth-limit
    while n = NULL do
        O ← {(⟨⟩, 0)}
        n ← SEARCHDL(O, d_max)
        d_max ← d_max + 1
    return n
```

Warning: Increasing $d_{\max}$ can be done in different ways.

2.5.3 Cost-Limiting

Definition: Cost limit of $c_{\max}$ to any search algorithm by modifying `SEARCH(·)` as follows:

```
procedure SEARCHCL(O, c_max):
    if O = ∅ then
        return NULL                                   ▷ the search algorithm failed to find a path to a goal
    n ← REMOVE(O)                                     ▷ "explore" a node, n
    if dst(n) ∈ G then
        return n                                      ▷ the search algorithm found a path to a goal
    for n' ∈ chl(n) do                                ▷ "expand" n and "export" its children
        if cst(n') ≤ c_max then                       ▷ unless the child is too expensive
            O ← O ∪ {n'}
    SEARCHCL(O, c_max)
```

2.5.4 Iterative-Inflating

Definition: Iteratively increase the cost limit, $c_{\max}$, to any search algorithm with cost-limiting, by placing `SEARCHCL(·)` in a wrapper, `SEARCHII(·)`:

```
procedure SEARCHII():
    n ← NULL
    c_max ← 0
    ▷ while a solution has not been found, reset the open set, run the search algorithm, then increase the
    cost-limit
    while n = NULL do
        O ← {(⟨⟩, 0)}
        n ← SEARCHCL(O, c_max)
        c_max ← c_max + ϵ
    return n
```

Warning: Increasing $c_{\max}$ can be done in different ways.

2.5.5 Intra-Path Cycle Checking

Definition: Do not expand a path if it is cyclic. Modify `SEARCH(·)` as follows:

```
procedure SEARCH(O):
    if O = ∅ then
        return NULL
    n ← REMOVE(O)
    if dst(n) ∈ G then
        return n
    for n' ∈ chl(n) do                                  ▷ "expand" n and "export" its children
        if not CYCLIC(n') then                                  ▷ unless the child is cyclic
            O ← O ∪ {n'}
    SEARCH(O)
```

- Optimately of an algorithm is preserved provided $\epsilon > 0$.

2.5.6 Inter-Path Cycle Checking

Definition: We modify `SEARCH(·)` as follows:

```
procedure SEARCH(O, C):
    if O = ∅ then
        return NULL
    n ← REMOVE(O)
    C ← C ∪ {n}                                                  ▷ add n to the closed set
    if dst(n) ∈ G then
        return n
    for n' ∈ chl(n) do                                  ▷ "expand" n and "export" its children
        if n' ∉ C then                                ▷ unless the child's destination is closed
            O ← O ∪ {n'}
    SEARCH(O, C)
```

and then call the algorithm as follows:

```
O ← {(⟨⟩, 0)}
C ← ∅                                                  ▷ initialize a set of closed vertices
SEARCH(O, C)
```

2.6 Informed Search Algorithms

2.6.1 Estimated Cost

Definition: ecst$(\cdot)$, to estimate the total cost to a goal given a path, p, based on the following:

- Cost of path p: cst(p)
- Estimate of the extra cost needed to get to a goal from dst(p): hur : $S \to \mathbb{R}_+$
  - hur(s) estimates the cost to get to $\mathcal{G}$ from s and hur(p) means hur(dst(p)).

Example: Some common choices for ecst$(\cdot)$ include:

1. ecst(p) = hur(p); called nearest-first search (NFS)
2. ecst(p) = cst(p) + hur(p); called A* (A-star)

2.7 Characteristics of an Informed Search Algorithm

Definition:

1. Heuristic: hur$(\cdot)$
2. Cost estimation: ecst$(\cdot)$

2.7.1 Heuristics
2.7.2 Heuristic-First Search (HFS)
2.7.3 A-Star Search (A*)
2.7.4 Iterative Inflating A-Star Search (IIA*)
2.7.5 Designing Heuristics via Problem Relaxation
2.7.6 Combining Heuristics
2.8 Anytime Search Algorithms
2.9 Formulating a Search Problem
3 Informed Search Algorithms
4 Probability Review
5 Probabilistic Inference Problems